

Virtualization Guide

5.1

Red Hat Enterprise Linux



Virtualization_Guide

ISBN: N/A

Publication date: January 2008

This Guide contains information on configuring, creating and monitoring guest operating systems on Red Hat Enterprise Linux 5.1, using virsh, vmm, and xend. If you find an error in the Red Hat Enterprise Linux Virtualization Guide, or if you have thought of a way to make this manual better, we would like to hear from you! Submit a report in Bugzilla (<http://bugzilla.redhat.com/bugzilla/>) against the product Red Hat Enterprise Linux and the component Virtualization_Guide.

Virtualization Guide: Red Hat Enterprise Linux

Copyright © 2008 Red Hat, Inc.

Copyright © 2008 Red Hat, Inc. This material may only be distributed subject to the terms and conditions set forth in the Open Publication License, V1.0 or later with the restrictions noted below (the latest version of the OPL is presently available at <http://www.opencontent.org/openpub/>).

Distribution of substantively modified versions of this document is prohibited without the explicit permission of the copyright holder.

Distribution of the work or derivative of the work in any standard (paper) book form for commercial purposes is prohibited unless prior permission is obtained from the copyright holder.

Red Hat and the Red Hat "Shadow Man" logo are registered trademarks of Red Hat, Inc. in the United States and other countries.

All other trademarks referenced herein are the property of their respective owners.

The GPG fingerprint of the security@redhat.com key is:

CA 20 86 86 2B D6 9D FC 65 F6 EC C4 21 91 80 CD DB 42 A6 0E

1801 Varsity Drive
Raleigh, NC 27606-2072
USA
Phone: +1 919 754 3700
Phone: 888 733 4281
Fax: +1 919 754 3701
PO Box 13588
Research Triangle Park, NC 27709
USA

1. Red Hat Virtualization System Architecture	1
2. Operating System Support	3
3. Hardware Support	5
4. Red Hat Virtualization System Requirements	7
5. Booting the System	9
6. Configuring GRUB	11
7. Booting a Guest Domain	13
8. Starting/Stopping a Domain at Boot Time	15
9. Configuration Files	17
10. Managing CPUs	19
11. Migrating a Domain	21
12. Configuring for Use on a Network	23
13. Securing Domain0	25
14. Storage	27
15. Managing Virtual Machines with virsh	29
1. Connecting to a Hypervisor	29
2. Creating a Virtual Machine	29
3. Configuring an XML Dump	29
4. Suspending a Virtual Machine	29
5. Resuming a Virtual Machine	30
6. Saving a Virtual Machine	30
7. Restoring a Virtual Machine	30
8. Shutting Down a Virtual Machine	30
9. Rebooting a Virtual Machine	31
10. Terminating a Domain	31
11. Converting a Domain Name to a Domain ID	31
12. Converting a Domain ID to a Domain Name	31
13. Converting a Domain Name to a UUID	31
14. Displaying Virtual Machine Information	32
15. Displaying Node Information	32
16. Displaying the Virtual Machines	32
17. Displaying Virtual CPU Information	33
18. Configuring Virtual CPU Affinity	33
19. Configuring Virtual CPU Count	33
20. Configuring Memory Allocation	34
21. Configuring Maximum Memory	34
22. Managing Virtual Networks	34
16. Managing Virtual Machines Using xend	37
17. Managing Virtual Machines with Virtual Machine Manager	41
1. Virtual Machine Manager Architecture	41
2. The Open Connection Window	41
3. Virtual Machine Manager Window	42
4. Virtual Machine Details Window	43
5. Virtual Machine Graphical Console	43
6. Starting the Virtual Machine Manager	44
7. Creating a New Virtual Machine	45

8. Restoring A Saved Machine	54
9. Displaying Virtual Machine Details	56
10. Configuring Status Monitoring	61
11. Displaying Domain ID	62
12. Displaying Virtual Machine Status	64
13. Displaying Virtual CPUs	65
14. Displaying CPU Usage	67
15. Displaying Memory Usage	68
16. Managing a Virtual Network	70
17. Creating a Virtual Network	71
18. Red Hat Virtualization Troubleshooting	81
1. Logfile Overview and Locations	81
2. Logfile Descriptions	81
3. Important Directory Locations	82
4. Troubleshooting Tools	82
5. Troubleshooting with the Logs	84
6. Troubleshooting with the Serial Console	85
7. Paravirtualized Guest Console Access	85
8. Full Virtualization Guest Console Access	86
9. Implementing Lun Persistence	86
10. SELinux Considerations	88
11. Accessing Data on Guest Disk Image	88
12. Common Troubleshooting Situations	89
13. Loop Device Errors	90
14. Guest Creation Errors	90
15. Serial Console Errors	90
16. Network Bridge Errors	91
17. Laptop Configurations	92
18. Starting Domains Automatically During System Boot	94
19. Modifying Domain0	95
20. Guest Configuration Files	96
21. Cloning the Guest Configuration Files	97
22. Creating a Script to Generate MAC Addresses	97
23. Configuring Virtual Machine Live Migration	97
24. Interpreting Error Messages	98
25. Online Troubleshooting Resources	101
19. Additional Resources	103
1. Useful Websites	103
2. Installed Documentation	103
A. Lab 1	105
B. Lab 2	111

Red Hat Virtualization System Architecture

A functional Red Hat Virtualization system is multi-layered and is driven by the privileged Red Hat Virtualization component. Red Hat Virtualization can host multiple guest operating systems. Each guest operating system runs in its own domain, Red Hat Virtualization schedules virtual CPUs within the virtual machines to make the best use of the available physical CPUs. Each guest operating system handles its own applications. These guest operating systems schedule each application accordingly.

You can deploy Red Hat Virtualization in one of two choices: **full virtualization** or **paravirtualization**. Full virtualization provides total abstraction of the underlying physical system and creates a new virtual system in which the guest operating systems can run. No modifications are needed in the guest OS or application (the guest OS or application is not aware of the virtualized environment and runs normally). Paravirtualization requires user modification of the guest operating systems that run on the virtual machines (these guest operating systems are aware that they are running on a virtual machine) and provide near-native performance. You can deploy both paravirtualization and full virtualization across your virtualization infrastructure.

The first domain, known as **domain0** (dom0), is automatically created when you boot the system. Domain0 is the privileged guest and it possesses management capabilities which can create new domains and manage their virtual devices. Domain0 handles the physical hardware, such as network cards and hard disk controllers. Domain0 also handles administrative tasks such as suspending, resuming, or migrating guest domains to other virtual machines.

The **hypervisor** (Red Hat's Virtual Machine Monitor) is a virtualization platform that allows multiple operating systems to run on a single host simultaneously within a full virtualization environment. A guest is an operating system (OS) that runs on a virtual machine in addition to the host or main OS.

With Red Hat Virtualization, each guest's **memory** comes from a slice of the host's physical memory. For paravirtual guests, you can set both the initial memory and the maximum size of the virtual machine. You can add (or remove) physical memory to the virtual machine at runtime without exceeding the maximum size you specify. This process is called ballooning.

You can configure each guest with a number of virtual **cpus** (called vcpus). The Virtual Machine Manager schedules the vcpus according to the workload on the physical CPUs.

You can grant a guest any number of **virtual disks**. The guest sees these as either hard disks or (for full virtual guests) as CD-ROM drives. Each virtual disk is served to the guest from a block device or from a regular file on the host. The device on the host contains the entire full disk image for the guest, and usually includes partition tables, multiple partitions, and potentially LVM physical volumes.

Virtual networking interfaces runs on the guest. Other interfaces can run on the guest like

virtual ethernet internet cards (VNICs). These network interfaces are configured with a persistent virtual media access control (MAC) address. The default installation of a new guest installs the VNIC with a MAC address selected at random from a reserved pool of over 16 million addresses, so it is unlikely that any two guests will receive the same MAC address. Complex sites with a large number of guests can allocate MAC addresses manually to ensure that they remain unique on the network.

Each guest has a virtual **text console** that connects to the host. You can redirect guest logins and console output to the text console.

You can configure any guest to use a virtual **graphical console** that corresponds to the normal video console on the physical host. You can do this for full virtual and paravirtual guests. It employs the features of the standard graphic adapter like boot messaging, graphical booting, multiple virtual terminals, and can launch the x window system. You can also use the graphical keyboard to configure the virtual keyboard and mouse.

Guests can be identified in any of three **identities**: domain name (domain-name), identity (domain-id), or UUID. The domain-name is a text string that corresponds to a guest configuration file. The domain-name is used to launch the guests, and when the guest runs the same name is used to identify and control it. The domain-id is a unique, non-persistent number that gets assigned to an active domain and is used to identify and control it. The UUID is a persistent, unique identifier that is controlled from the guest's configuration file and ensures that the guest is identified over time by system management tools. It is visible to the guest when it runs. A new UUID is automatically assigned to each guest by the system tools when the guest first installs.

Operating System Support

Red Hat Virtualization's paravirtualization mode allows you to utilize high performance virtualization on architectures that are potentially difficult to virtualize such as x86 based systems. To deploy para-virtualization across your operating system(s), you need access to the paravirtual guest kernels that are available from a respective Red Hat distro (for example, RHEL 4.0, RHEL 5.0, etc.). Whilst your operating system kernels must support Red Hat Virtualization, it is not necessary to modify user applications or libraries.

Red Hat Virtualization allows you to run an unmodified guest kernel if you have Intel VT and AMD SVM CPU hardware. You do not have to port your operating system to deploy this architecture on your Intel VT or AMD SVM systems. Red Hat Virtualization supports:

- Intel VT-x or AMD-V Pacifica and Vanderpool technology for full and paravirtualization.
- Intel VT-i for ia64
- Linux and UNIX operating systems, including NetBSD, FreeBSD, and Solaris.
- Microsoft Windows as an unmodified guest operating system with Intel Vanderpool or AMD's Pacifica technology.

To run full virtualization guests on systems with Hardware-assisted Virtual Machine (HVM), Intel, or AMD platforms, you must check to ensure your CPUs have the capabilities needed to do so.

To check if you have the CPU flags for Intel support, enter the following:

```
grep vmx /proc/cpuinfo
```

The output displays:

```
flags      :  fpu tsc msr pae mce cx8 apic mtrr mca cmov pat pse36 clflush dts
acpi mmx fxsr sse sse2 ss ht tm syscall nx lm constant_tsc pni monitor
ds_cpl vmx est tm2 cx16 xtpr lahf_lm
```

If a vmx flag appears then you have Intel support.

To check if you have the CPU flags for AMD support, enter the following:

```
grep svm /proc/cpuinfo
cat /proc/cpuinfo | grep svm
```

The output displays:

```
flags      :  fpu tsc msr pae mce cx8 apic mtrr mca cmov pat pse36 clflush dt
acpi mmx fxsr sse sse2 ss ht tm syscall nx mmtxt fxsr_opt rdtscp lm
3dnowext  pni  cx16  lahf_lm  cmp_legacy  svm  cr8_legacy
```

If an svm flag appears then you have AMD support.



note

In addition to checking the CPU flags, you should enable full virtualization within your system BIOS.

Hardware Support

Red Hat Virtualization supports multiprocessor systems and allows you to run Red Hat Virtualization on x86 architected systems with a P6 class (or earlier) processors like:

- Celeron
- Pentium II
- Pentium III
- Pentium IV
- Xeon
- AMD Athlon
- AMD Duron

With Red Hat Virtualization, 32-bit hosts runs only 32-bit paravirtual guests. 64-bit hosts runs only 64-bit paravirtual guests. And a 64-bit full virtualization host runs 32-bit, 32-bit PAE, or 64-bit guests. A 32-bit full virtualization host runs both PAE and non-PAE full virtualization guests.

The Red Hat Enterprise Linux Virtualization kernel does not support more than 32GB of memory for x86_64 systems. If you need to boot the virtualization kernel on systems with more than 32GB of physical memory installed, you must append the kernel command line with `mem=32G`. This example shows how to enable the proper parameters in the `grub.conf` file:

```
title Red Hat Enterprise Linux Server (2.6.18-4.elxen)
root (hd0, 0)
kernel /xen.gz-2.6.18-4-el5 mem=32G
module /vmlinuz -2.6.18-4.el5xen ro root=LABEL=/
module /initrd-2.6.18-4.el5xen.img
```

PAE (Physical Address Extension) is a technology that increases the amount of physical or virtual memory available to user applications. Red Hat Virtualization requires that PAE is active on your systems. Red Hat Virtualization 32 bit architecture with PAE supports up to 16 GB of physical memory. It is recommended that you have at least 256 megabytes of RAM for every guest you have running on the system. Red Hat Virtualization enables x86/64 machines to address up to physical 64 GB. The Red Hat Virtualization kernels will not run on a non-PAE system. To determine if a system supports PAE, type the following commands:

```
grep pae /proc/cpuinfo
```

The following output displays:

```
flags : fpu tsc msr pae mce cx8 apic mtrr mca cmov pat pse36 mmx fxsr sse  
syscall mmtxt 3dnowext 3dnow up ts
```

If your output matches (or is similar to) the above, then your CPU supports PAE. If the command prompt displays nothing, then your CPU does not support PAE.

Red Hat Virtualization System Requirements

The items listed below are required by the Red Hat Virtualization system:

- A working Red Hat RHEL 5 Linux distribution
- A working GRUB bootloader
- Root access
- A P6 class (or earlier) processor
- The Linux bridge-utils
- The Linux hotplug systems
- zlib development installation
- Python 2.2 runtime
- initscripts

The dependencies are configured automatically during the installation process.



Note

If your system CPU architecture is ia64, you need to manually install the `xen-ia64-guest-firmware` package to run a fully virtualized guest. This package is provided in the Supplementary CD and is not installed by default.

Booting the System

After installing the Red Hat Virtualization components, you must reboot the system. When the boot completes, you must log into your system as usual. Then before you start Red Hat Virtualization you must log in a root. The xend control daemon should already be initiated by initscripts, but to start the xend manually, enter:

```
service xend start
```

You can also use `chkconfig xend` when installing to enable `xend` at boot time.

The xend node control daemon performs system management functions that relate to virtual machines. This daemon controls the virtualized resources, and xend must be running to interact with virtual machines. Before you start xend, you must specify the operating parameters by editing the xend configuration file `xend-config.sxp` which is located in the `etc/xen` directory.

Configuring GRUB

GNU Grand Unified Boot Loader (or GRUB) is a program which enables the user to select which installed operating system or kernel to load at system boot time. It also allows the user to pass arguments to the kernel. The GRUB configuration file (located in `/boot/grub/grub.conf`) is used to create a list of operating systems to boot in GRUB's menu interface. When you install the kernel-xen RPM, a post script adds kernel-xen entries to the GRUB configuration file. You can edit the `grub.conf` file and enable the following GRUB parameter:

```
title Red Hat Enterprise Linux Server (2.6.18-3.el5xen)
root    (hd0; 0)
kernel  /xen.gz.-2.6.18-3.el5
module  /vmlinuz-2.6..18-3.el5xen ro root=/dev/VolGroup00/LogVol100 rhgb
quiet
module  /initrd-2.6.18-3. el5xenxen.img
```

If you set your Linux grub entries to reflect this example, the boot loader loads the hypervisor, `initrd` image, and Linux kernel. Since the kernel entry is on top of the other entries, the kernel loads into memory first. The boot loader sends (and receives) command line arguments to and from the hypervisor and Linux kernel. This example entry shows how you would restrict the Domain0 linux kernel memory to 800 MB:

```
title Red Hat Enterprise Linux Server (2.6.18-3.el5xen)
root    (hd0; 0)
kernel  /xen.gz.-2.6.18-3.el5 dom0_mem=800M
module  /vmlinuz-2.6..18-3.el5xen ro root=/dev/VolGroup00/LogVol100 rhgb
quiet
module  /initrd-2.6.18-3. el5xenxen.img
```

You can use these GRUB parameters to configure the Virtualization hypervisor:

```
mem
```

This limits the amount of memory that is available for domain0.

```
com1=115200, 8n1
```

This enables the first serial port in the system to act as serial console (com2 is assigned for the next port, and so on...).

```
dom0_mem
```

This limits the amount of memory that is available for domain0.

```
dom0_max_vcpus
```

This limits the amount of CPUs visible to domain0.

```
acpi
```

This switches the ACPI hypervisor to the hypervisor and domain0. The ACPI parameter options include:

```
/* **** Linux config options: propagated to domain0 ****/
/* "acpi=off":      Disables both ACPI table parsing and interpreter.    */
/* "acpi=force":    Overrides the disable blacklist.                      */
/* "acpi=strict":   Disables out-of-spec workarounds.                     */
/* "acpi=ht":       Limits ACPI from boot-time to enable HT.              */
/* "acpi=noirq":    Disables ACPI interrupt routing.                      */
```

```
noacpi
```

This disables ACPI for interrupt delivery.

Booting a Guest Domain

You can boot guest domains by using the `xm` application. You can also use `virsh` and the Virtual Machine Manager to boot the guests. A prerequisite for booting a guest domain is to install a guest host first. This example uses the `xm create` subcommand:

```
# xm create -c guestdomain1
```

The `guestdomain1` is the configuration file for the domain you are booting. The `-c` option connects to the actual console after booting.

Starting/Stopping a Domain at Boot Time

You can start or stop running domains at any time. Domain0 waits for all running domains to shutdown before restarting. You must place the configuration files of the domains you wish to shut down in the `/etc/xen/` directory. All the domains that you want to start at boot time must be symlinked to `/etc/xen/auto`.

```
chkconfig xendomains on
```

The `chkconfig xendomains on` command does not automatically start domains; instead it will start the domains on the next boot.

```
chkconfig xendomains off
```

Terminates all running Red Hat Virtualization domains. The `chkconfig xendomains off` command shuts down the domains on the next boot.

Configuration Files

Red Hat Virtualization configuration files contain the following standard variables. Configuration items within these files must be enclosed in quotes ("). These configuration files reside in the `/etc/xen` directory.

Item	Description
pae	Specifies the physical address extension configuration data.
apic	Specifies the advanced programmable interrupt controller configuration data.
memory	Specifies the memory size in megabytes.
vcpus	Specifies the numbers of virtual CPUs.
console	Specifies the port numbers to export the domain consoles to.
nic	Specifies the number of virtual network interfaces.
vif	Lists the randomly-assigned MAC addresses and bridges assigned to use for the domain's network addresses.
disk	Lists the block devices to export to the domain and exports physical devices to domain with read only access.
dhcp	Enables networking using DHCP.
netmask	Specifies the configured IP netmasks.
gateway	Specifies the configured IP gateways.
acpi	Specifies the advanced configuration power interface configuration data.

Table 9.1. Red Hat Virtualization Configuration Files

Managing CPUs

Red Hat Virtualization allows a domain's virtual CPUs to associate with one or more host CPUs. This can be used to allocate real resources among one or more guests. This approach allows Red Hat Virtualization to make optimal use of processor resources when employing dual-core, hyperthreading, or other advanced CPU technologies. If you are running I/O intensive tasks, its typically better to dedicate either a hyperthread or entire core to run domain0. The Red Hat Virtualization credit scheduler automatically rebalances virtual cpus between physical ones, to maximize system use. The Red Hat Virtualization system allows the credit scheduler to move CPUs around as necessary, as long as the virtual CPU is pinned to a physical CPU.

Migrating a Domain

Migration is the transferal of a running virtual domain from one physical host to another. Red Hat Virtualization supports two varieties of migration — offline and live. Offline migration moves a virtual machine from one host to another by pausing it, transferring its memory, and then resuming it on the host destination. Live migration does the same thing, but does not directly affect the domain. When performing a live migration, the domain continues its usual activities, and from the user perspective is unnoticeable. To initiate a live migration, both hosts must be running Red Hat Virtualization and the xend daemon. The destination host must have sufficient resources (such as memory capacity) to accommodate the domain bandwidth after the migration. Both the source and destination machines must have the same architecture and virtualization extensions (such as i386-VT, x86-64-VT, x86-64-SVM, etc.) and must be on the same L2 subnet.

When a domain migrates its MAC and IP addresses move with it. Only virtual machines with the same layer-2 network and subnets will successfully migrate. If the destination node is on a different subnet, the administrator must manually configure a suitable EtherIP or IP tunnel in the remote node of domain0. The xend daemon stops the domain and copies the job over to the new node and restarts it. The Red Hat Virtualization RPM does not enable migration from any other host except the localhost (see the `/etc/xend-config.sxp` file for information). To allow the migration target to accept incoming migration requests from remote hosts, you must modify the target's `xen-relocation-hosts-allow` parameter. Be sure to carefully restrict which hosts are allowed to migrate, since there is no authentication.

Since these domains have such large file allocations, this process can be time consuming. If you migrate a domain with open network connections, they will be preserved on the host destination, and SSH connections should still function. The default Red Hat Virtualization iptables rules will not permit incoming migration connections. To allow this, you must create explicit iptables rules.

You can use the `xm migrate` command to perform an offline migration :

```
xm migrate domain-id [destination domain]
```

You can use the `xm migrate` command to perform a live migration:

```
xm migrate domain-id -l [destination domain]
```

You may need to reconnect to the domain's console on the new machine. You can use the `xm console` command to reconnect.

Configuring for Use on a Network

Integrating Red Hat Virtualization into your network architecture is a complicated process and depending upon your infrastructure, may require custom configuration to deploy multiple ethernet interfaces and setup bridging.

Each domain network interface is connected to a virtual network interface in `dom0` by a point to point link. These devices are `vif<domid>` and `<vifid>`. `vif1.0` for the first interface in domain 1; `vif3.1` for the second interface in domain 3.

Domain0 handles traffic on these virtual interfaces by using standard Linux conventions for bridging, routing, rate limiting, etc. The **xend** daemon employs two shell scripts to perform initial configuration of your network and new virtual interfaces. These scripts configure a single bridge for all virtual interfaces. You can configure additional routing and bridging by customizing these scripts.

Red Hat Virtualization's virtual networking is controlled by the two shell scripts, `network-bridge` and `vif-bridge`. **xend** calls these scripts when certain events occur. Arguments can be passed to the scripts to provide additional contextual information. These scripts are located in the `/etc/xen/scripts` directory. You can change script properties by modifying the `xend-config.sxp` configuration file located in the `/etc/xen` directory.

`network-bridge` — When **xend** is started or stopped, this script initializes or shuts down the virtual network. Then the configuration initialization creates the bridge `xen-br0` and moves `eth0` onto that bridge, modifying the routing accordingly. When **xend** finally exits, it deletes the bridge and removes `eth0`, thereby restoring the original IP and routing configuration.

`vif-bridge` is a script that is invoked for every virtual interface on the domain. It configures firewall rules and can add the `vif` to the appropriate bridge.

There are other scripts that you can use to help in setting up Red Hat Virtualization to run on your network, such as `network-route`, `network-nat`, `vif-route`, and `vif-nat`. Or these scripts can be replaced with customized variants.

Securing Domain0

When deploying Red Hat Virtualization on your corporate infrastructure, you must ensure that domain0 cannot be compromised. Domain0 is the privileged domain that handles system management. If domain0 is insecure, all other domains in the system are vulnerable. There are several ways to implement security you should know about when integrating Red Hat Virtualization into your systems. Together with other people in your organization, you should create a 'deployment plan' that contains the operating specifications and services that will run on Red Hat Virtualization, and what is needed to support these services. Here are some security issues to consider when putting together a deployment plan:

- Run the lowest number of necessary services. You do not want to include too many jobs and services in domain0. The less things running on domain0, the higher the level of security.
- Enable SELINUX to help secure domain0.
- Use a firewall to restrict traffic to domain0. You can setup a firewall with default-reject rules that will help secure attacks on domain0. It is also important to limit network facing services.
- Do not allow normal users to access domain0. If you do permit normal users domain0 access, you run the risk of rendering domain0 vulnerable. Remember, domain0 is privileged, and granting unprivileged accounts may compromise the level of security.

Storage

There are several ways to manage virtual machine storage. You can export a domain0 physical block device (hard drive or partition) to a guest domain as a virtual block device (VBD). You can also export directly from a partitioned image as a file-backed VBD. Red Hat Virtualization enables LVM and blktap by default during installation. You can also employ standard network protocols such as NFS, CLVM, or iSCSI to provide storage for virtual machines.

Managing Virtual Machines with **virsh**

You can use the **virsh** application to manage virtual machines. This utility is built around the libvirt management API and operates as an alternative to the **xm** tool or the graphical Virtual Machine Manager. Unprivileged users can employ this utility for read-only operations. If you plan on running `xend/qemu`, you should enable `xend/qemu` to run as a service. After modifying the respective configuration file, reboot the system, and `xend/qemu` will run as a service. You can use **virsh** to script vm work. Like the **xm** tool, you run **virsh** from the command line.

1. Connecting to a Hypervisor

You can use **virsh** to initiate a hypervisor session:

```
virsh connect <name>
```

Where `<name>` is the machine name of the hypervisor. If you want to initiate a read-only connection, append the above command with `--readonly`.

2. Creating a Virtual Machine

You can make a new virtual machine session from an XML machine definition. If you have a pre-existing guest that you created previously with the **xm** tool, you can also create a virtual machine for it:

```
virsh create <path to XML configuration file>
```

3. Configuring an XML Dump

You can use **virsh** to perform a data dump for an existing virtual machine.

```
virsh dumpxml [domain-id | domain-name | domain-uuid]
```

This command outputs the domain information (in XML) to `stdout`. If you save the data to a file, you can use the [create option](#) to recreate the virtual machine.

4. Suspending a Virtual Machine

You can use **virsh** to suspend a domain:

```
virsh suspend [domain-id | domain-name | domain-uuid]
```

When a domain is in a suspended state, it still consumes system RAM. There will also be no disk or network I/O when suspended. This operation is immediate and the virtual machine must be restarted with the [resume option](#).

5. Resuming a Virtual Machine

You can use **virsh** to restore a suspended virtual machine:

```
virsh resume [domain-id | domain-name | domain-uuid]
```

This operation is immediate and the virtual machine parameters are preserved in a `suspend` and `resume` cycle.

6. Saving a Virtual Machine

You can use **virsh** to save the current state of a virtual machine to a file:

```
virsh save [domain-name][domain-id | domain-uuid][filename]
```

This stops the virtual machine you specify and saves the data to a file, which may take some time given the amount of memory in use by your virtual machine. You can restore the state of the virtual machine with the [restore option](#).

7. Restoring a Virtual Machine

You can use **virsh** to restore a virtual machine that you previously saved with the [virsh save option](#):

```
virsh restore [filename]
```

This restarts the saved virtual machine, which may take some time. The virtual machine's name and UUID are preserved but are allocated for a new id.

8. Shutting Down a Virtual Machine

You can use **virsh** to shut down a virtual machine:

```
virsh shutdown [domain-id | domain-name | domain-uuid]
```

You can control the behavior of the rebooting virtual machine by modifying the `on_shutdown` parameter of the `xmdomain.cfg` file.

9. Rebooting a Virtual Machine

You can use **virsh** to reboot a virtual machine:

```
virsh reboot [domain-id | domain-name | domain-uuid]
```

You can control the behavior of the rebooting virtual machine by modifying the `on_reboot` parameter of the `xmdomain.cfg` file.

10. Terminating a Domain

You can use **virsh** to terminate a virtual machine:

```
virsh destroy [domain-name | domain-id | domain-uuid]
```

This command does an immediate ungraceful shutdown and stops any guest domain sessions (which could potentially lead to file corrupted filesystems still in use by the virtual machine). You should use the `destroy` option only when the virtual machine's operating system is non-responsive. For a paravirtualized virtual machine, you should use the [shutdown option](#).

11. Converting a Domain Name to a Domain ID

You can use **virsh** to convert a domain name or UUID to a domain id:

```
virsh domid [domain-name | domain-uuid]
```

12. Converting a Domain ID to a Domain Name

You can use **virsh** to convert a domain id or UUID to a domain name:

```
virsh domname [domain-name | domain-uuid]
```

13. Converting a Domain Name to a UUID

You can use **virsh** to convert a domain name to a UUID:

```
virsh domuuid [domain-id | domain-uuid]
```

14. Displaying Virtual Machine Information

You can use **virsh** to display information for a given virtual machine identified by its domain ID, domain name, or UUID:

```
virsh dominfo [domain-id | domain-name | domain-uuid]
```

15. Displaying Node Information

You can use **virsh** to display node information:

```
virsh nodeinfo
```

The outputs displays something similar to:

```
CPU model           x86_64
CPU (s)             8
CPU frequency       2895 Mhz
CPU socket(s)       2
Core(s) per socket  2
Threads per core:   2
Numa cell(s)        1
Memory size:        1046528 kb
```

This displays the node information and the machines that support the virtualization process.

16. Displaying the Virtual Machines

You can use **virsh** to display the virtual machine list and the current state:

```
virsh list domain-name [ --inactive | --all]
```

The `--inactive` option lists inactive domains (domains that have been defined but are not currently active). The `--all` domain lists all domains, whether active or not. Your output should resemble the this example:

ID	Name	State
0	Domain0	running
1	Domain202	paused
2	Domain010	inactive
3	Domain9600	crashed

Here are the six domain states:

```
running  lists domains currently active on the CPU
blocked  lists domains that are blocked
paused   lists domains that are suspended
shutdown lists domains that are in process of shutting down
shutoff  lists domains that are completely down.
crashed  lists domains that are crashed
```

17. Displaying Virtual CPU Information

You can use **virsh** to display virtual CPU information from a virtual machine:

```
virsh vcpuinfo [domain-id | domain-name | domain-uuid]
```

18. Configuring Virtual CPU Affinity

You can use **virsh** to configure the affinity of virtuals CPUs with physical CPUs:

```
virsh vcpupin [domain-id | domain-name | domain-uuid] [vcpu] , [cpulist]
```

Where `[vcpu]` is the virtual VCPU number and `[cpulist]` lists the physical number of CPUs.

19. Configuring Virtual CPU Count

You can use **virsh** to modify a Virtual Machine's number of CPUs:

```
virsh setvcpus [domain-name | domain-id | domain-uuid] [count]
```

Note that the new count cannot exceed the amount you specified when you created the Virtual Machine.

20. Configuring Memory Allocation

You can use **virsh** to modify a domain's memory allocation:

```
virsh setmem [domain-id | domain-name] [count]
```

You must specify the [count] in kilobytes. Note that the new count cannot exceed the amount you specified when you created the Virtual Machine. Values lower than 64 MB probably won't work. You can adjust the Virtual Machine memory as necessary.

21. Configuring Maximum Memory

You can use **virsh** to modify a Virtual Machine's maximum memory:

```
virsh setmaxmem [domain-name | domain-id | domain-uuid] [count]
```

You must specify the [count] in kilobytes. Note that the new count cannot exceed the amount you specified when you created the Virtual Machine. Values lower than 64 MB probably won't work. The maximum memory doesn't affect the current use of the Virtual Machine (unless the new value is lower which should shrink memory usage).

22. Managing Virtual Networks

You can use **virsh** to manage virtual networks. To list virtual networks:

```
virsh net-list
```

This command generates output similar to:

```
[root@domain ~]# virsh net-list
Name                State      Autostart
-----
default             active    yes
vnet1               active    yes
vnet2               active    yes
```

To view network information for a specific virtual network:


```
virsh net-dumpxml [vnet name]
```

This displays information about a specified virtual network in XML format:

```
[root@domain ~]# virsh net-dumpxml vnet1
<network>
  <name>vnet1</name>
  <uuid>98361b46-1581-acb7-1643-85a412626e70</uuid>
  <forward dev='eth0' />
  <bridge name='vnet0' stp='on' forwardDelay='0' />
  <ip address='192.168.100.1' netmask='255.255.255.0'>
    <dhcp>
      <range start='192.168.100.128' end='192.168.100.254' />
    </dhcp>
  </ip>
</network>
```

Other **virsh** commands used in managing virtual networks are:

- `virsh net-autostart [network name]` — Autostart a network specified as [network name]
- `virsh net-create [XML file]` — Generates and starts a new network using a preexisting XML file
- `virsh net-define [XML file]` — Generates a new network from a preexisting XML file without starting it
- `virsh net-destroy [network name]` — Destroy a network specified as [network name]
- `virsh net-name [network UUID]` — Convert a specified [network UUID] to a network name
- `virsh net-uuid [network name]` — Convert a specified [network name] to a network UUID
- `virsh net-start [name of an inactive network]` — Starts a previously undefined inactive network
- `virsh net-undefine [name of an inactive network]` — Undefine an inactive network

Managing Virtual Machines Using **xend**

The **xend** node control daemon performs certain system management functions that relate to virtual machines. This daemon controls the virtualized resources, and **xend** must be running to interact with virtual machines. Before you start **xend**, you must specify the operating parameters by editing the **xend** configuration file `xend-config.sxp` which is located in the `etc/xen` directory. Here are the parameters you can enable or disable in the `xend-config.sxp` configuration file:

Item	Description
console-limit	Determines the console server's memory buffer limit and assigns values on a per-domain basis
min-mem	Determines the minimum number of megabytes that is reserved for domain0 (if you enter 0, the value does not change)
dom0 cpus	Determines the number of CPUs in use by domain0 (at least 1 CPU is assigned by default)
enable-dump	Determines that a crash occurs then enables a dump (default is 0)
external-migration-tool	Determines the script or application that handles external device migration (scripts must reside in <code>etc/xen/scripts/external-device-migrate</code>)
logfile	Determines the location of the log file (default is <code>/var/log/xend.log</code>)
loglevel	Filters out the log mode values: DEBUG, INFO, WARNING, ERROR, or CRITICAL (default is DEBUG)
network-script	Determines the script that enables the networking environment (scripts must reside in <code>etc/xen/scripts</code> directory)
xend-http-server	Enables the http stream packet management

Item	Description
	server (default is no)
xend-unix-server	Enables the unix domain socket server (a socket server is a communications endpoint that handles low level network connections and accepts or rejects incoming connections)
xend-relocation-server	Enables the relocation server for cross-machine migrations (default is no)
xend-unix-path	Determines the location where the <code>xend-unix-server</code> command outputs data (default is <code>var/lib/xend/xend-socket</code>)
xend-port	Determines the port that the http management server uses (default is 8000)
xend-relocation-port	Determines the port that the relocation server uses (default is 8002)
xend-relocation-address	Determines the virtual machine addresses that are allowed for system migration
xend-address	Determines the address that the domain socket server binds to.

Table 16.1. Red Hat Virtualization xend Configuration Parameters

After setting these operating parameters, you should verify that xend is running and if not, initialize the daemon. At the command prompt, you can start the **xend** daemon by entering the following:

```
service xend start
```

You can use **xend** to stop the daemon:

```
service xend stop
```

This stops the daemon from running.

You can use **xend** to restart the daemon:

```
service xend restart
```

The daemon starts once again.

You check the status of the **xend** daemon.

```
service xend status
```

The output displays the daemon's status.

Managing Virtual Machines with Virtual Machine Manager

This section describes the Red Hat Virtualization Virtual Machine Manager (VMM) windows, dialog boxes, and various GUI controls.

1. Virtual Machine Manager Architecture

Red Hat Virtualization is a collection of software components that work together to host and manage virtual machines. The Virtual Machine Manager (VMM) gives you a graphical view of the virtual machines on your system. You can use VMM to define both para-virtual and full virtual machines. Using Virtual Machine Manager, you can perform any number of virtualization management tasks including assigning memory, assigning virtual CPUs, monitoring operational performance, and save, restore, pause, resume, and shutdown virtual systems. It also allows you to access the textual and graphical console. Red Hat Virtualization abstracts CPU and memory resources from the underlying hardware and network configurations. This enables processing resources to be pooled and dynamically assigned to applications and service requests. Chip-level virtualization enables operating systems with Intel VT and AMD Pacifica hardware to run on hypervisors.

2. The Open Connection Window

This window appears first and prompts the user to choose a hypervisor session. Non-privileged users can initiate a read-only session. Root users can start a session with full blown read-write status. For normal use, select the **Local Xen host** option. You start the Virtual Machine Manager test mode by selecting the **Other hypervisor** and then type `test:///default` in the URL field beneath. Once in test mode, you can connect to a libvirt dummy hypervisor. Note that although the **Remote Xen host screen** is visible, the functionality to connect to such a host is not implemented into Red Hat Enterprise Linux 5.1.



Figure 17.1. Virtual Machine Manager Connection window

3. Virtual Machine Manager Window

This main window displays all the running virtual machines and resources currently allocated to them (including domain0). You can decide which fields to display. Double-clicking on the desired virtual machine brings up the respective console for that particular machine. Selecting a virtual machine and double-click the **Details** button to display the Details window for that machine. You can also access the **File** menu to create a new virtual machine.

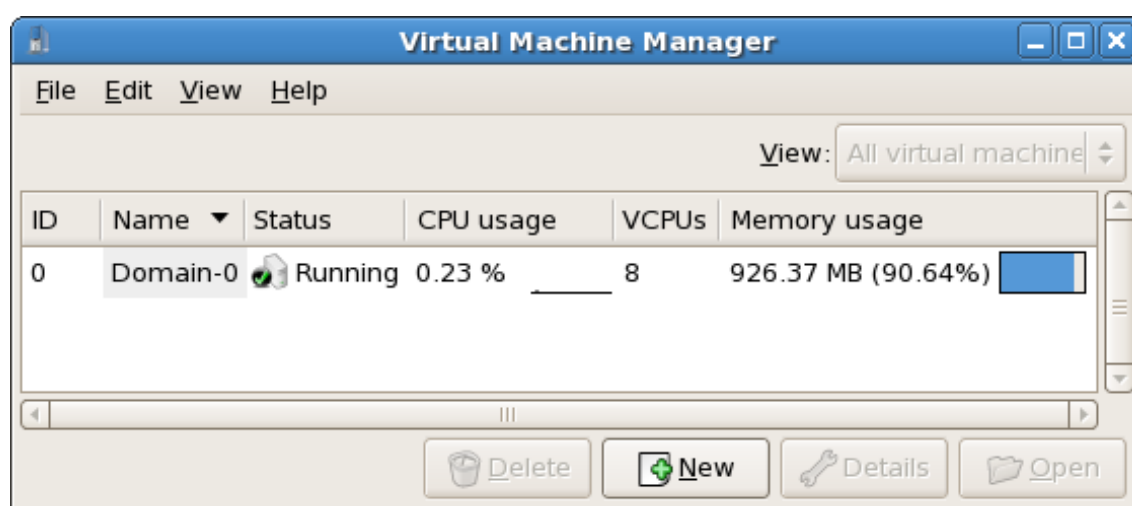


Figure 17.2. Virtual Machine Manager main window

4. Virtual Machine Details Window

This window displays graphs and statistics of a guest's live resource utilization data available from the Red Hat Virtualization Virtual Machine Manager. The UUID field displays the globally unique identifier for the virtual machines(s).




Figure 17.3. Virtual Machine Manager Details window

5. Virtual Machine Graphical Console

This window displays a virtual machine's graphical console. Paravirtual and full virtual machines use different techniques to export their local virtual framebuffers, but both technologies use VNC to make them available to the Virtual Machine Manager's console window. If your virtual machine is set to require authentication, the Virtual Machine Graphical console prompts you for a password before the display appears.




Figure 17.4. Graphical Console window

Your local desktop can intercept key combinations (for example, Ctrl+Alt+F11) to prevent them from being sent to the guest machine. You can use the Virtual Machine Manager's 'sticky key' capability to send these sequences. You must press any modifier key (like Ctrl or Alt) 3 times and the key you specify gets treated as active until the next non-modifier key is pressed. Then you can send Ctrl-Alt-F11 to the guest by entering the key sequence 'Ctrl Ctrl Ctrl Alt+F1'.

6. Starting the Virtual Machine Manager

To start the Virtual Machine Manager session, from the Applications menu, click System Tools and select Virtual Machine Manager .

The Virtual Machine Manager main window appears.




Figure 17.5. Starting the Virtual Machine Manager

7. Creating a New Virtual Machine

The **Virtual Machine Manager** (virt-manager) is the desktop application that manages virtual machines.

You can use Red Hat's Virtual Machine Manager to:

- Create new domains.
- Configure or adjust a domain's resource allocation and virtual hardware.
- Summarize running domains with live performance and resource utilization statistics.
- Display graphs that show performance and resource utilization over time.
- Use the embedded VNC client viewer which presents a full graphical console to the guest domain.



Note:

You must install Red Hat Enterprise Linux 5.1, virt-manager, and the kernel packages on all systems that require virtualization. All systems then must be booted and running the Red Hat Virtualization kernel.

These are the steps required to install a guest operating system on Red Hat Enterprise Linux 5 using the Virtual Machine Monitor:

Procedure 17.1. Creating a Guest Operating System

1. From the **Applications** menu, select **System Tools** and then **Virtual Machine Manager**.

The Virtual Machine Manager main window appears.




Figure 17.6. Virtual Machine Manager window

2. From the **File** menu, select **New machine**.

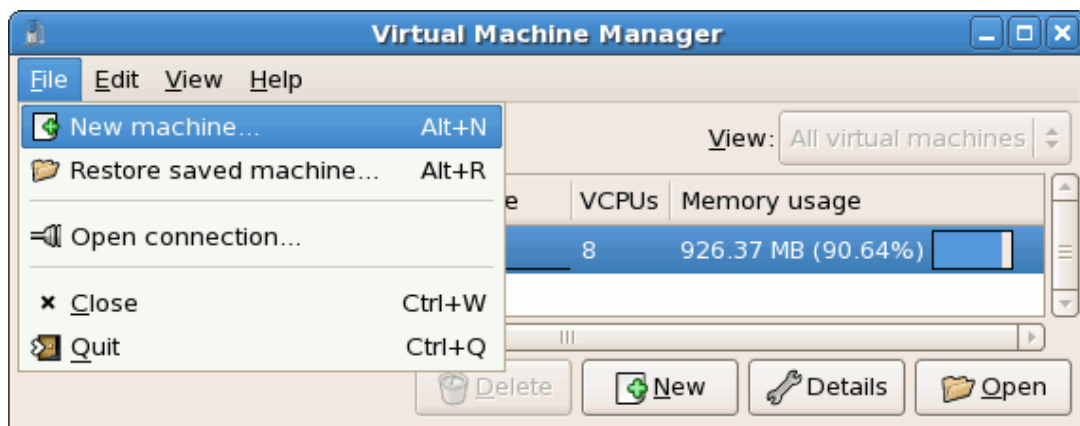


Figure 17.7. Selecting a New Machine

The Creating a new virtual system wizard appears.

3. Click **Forward**.

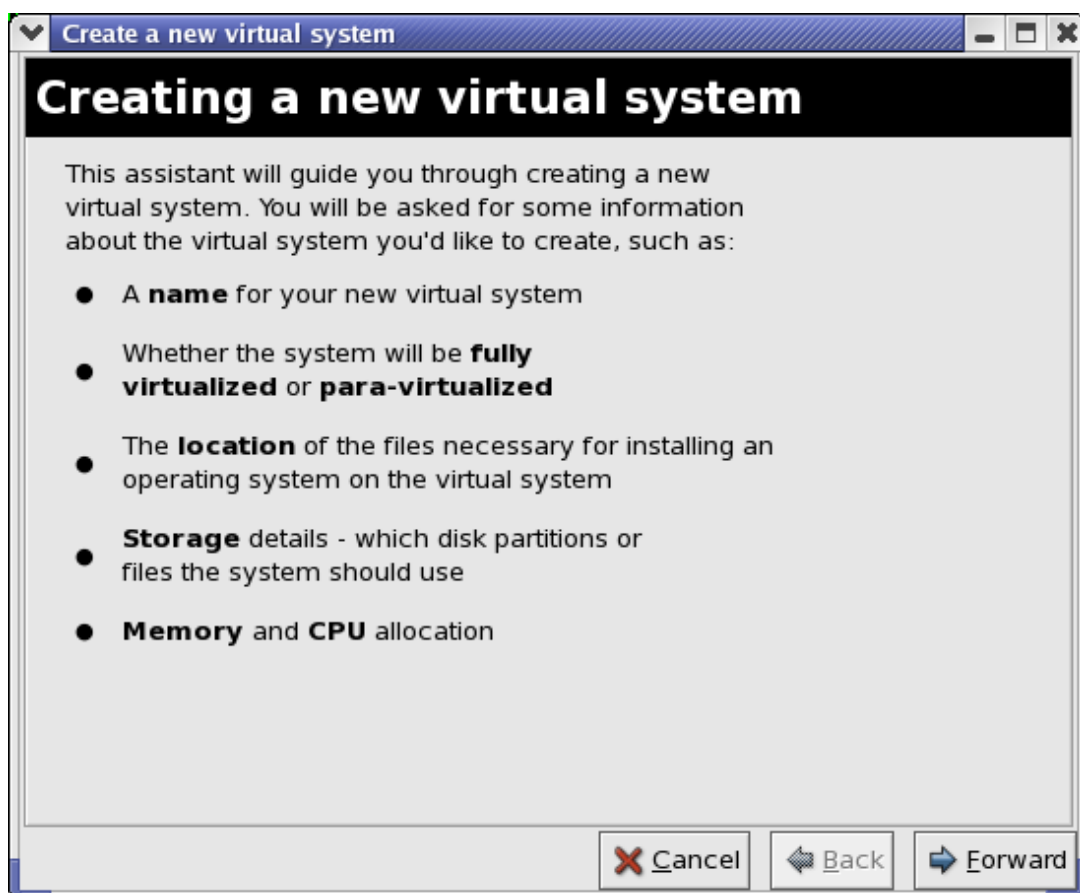


Figure 17.8. Creating a New Virtual System Wizard

4. Enter the name of the new virtual system and then click **Forward**.

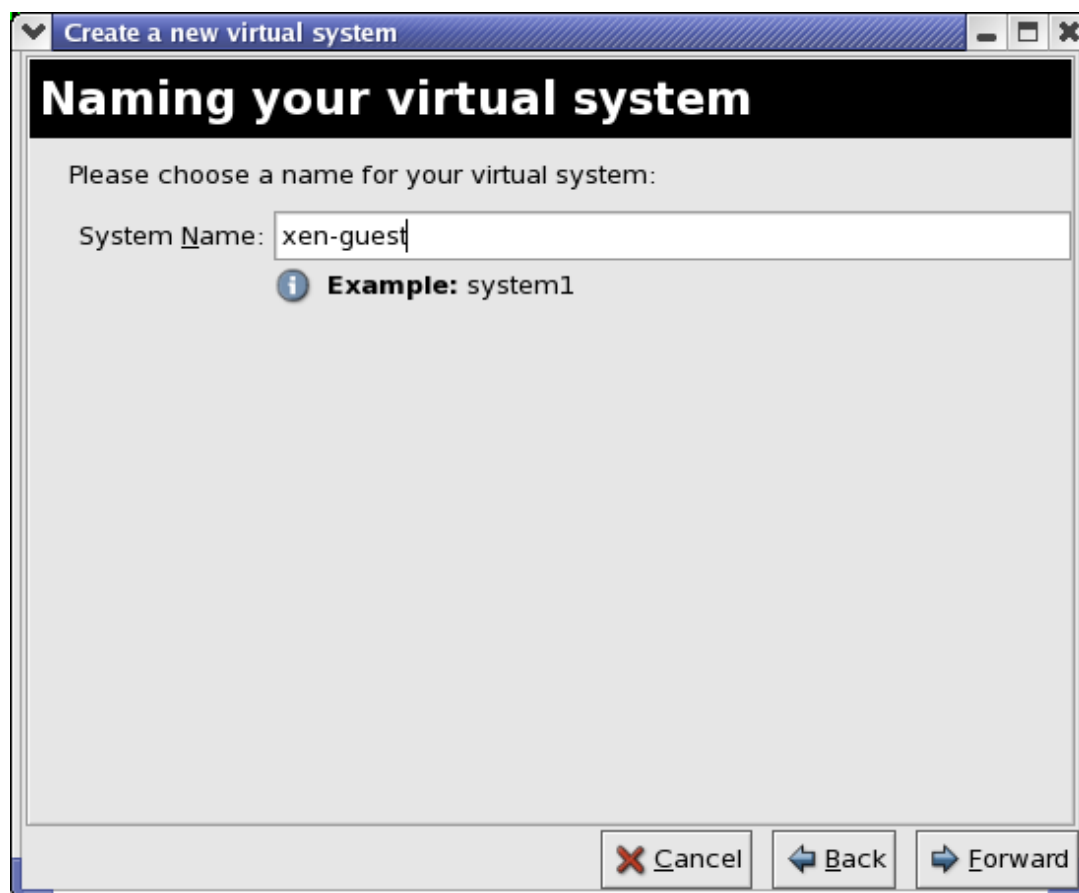


Figure 17.9. Naming the Virtual System

5. Enter the location of your install media. Location of the kickstart file is optional. Then click **Forward**.




Figure 17.10. Locating the Installation Media

6. Install either to a physical disk partition or install to a virtual file system within a file.



Note

This example installs a virtual system within a file.

SELinux policy only allows xen disk images to reside in `/var/lib/xen/images`.

Open a terminal and create the `/xen` directory and set the SELinux policy with the command `restorecon -v /xen`. Specify your location and the size of the virtual disk, then click **Forward**.

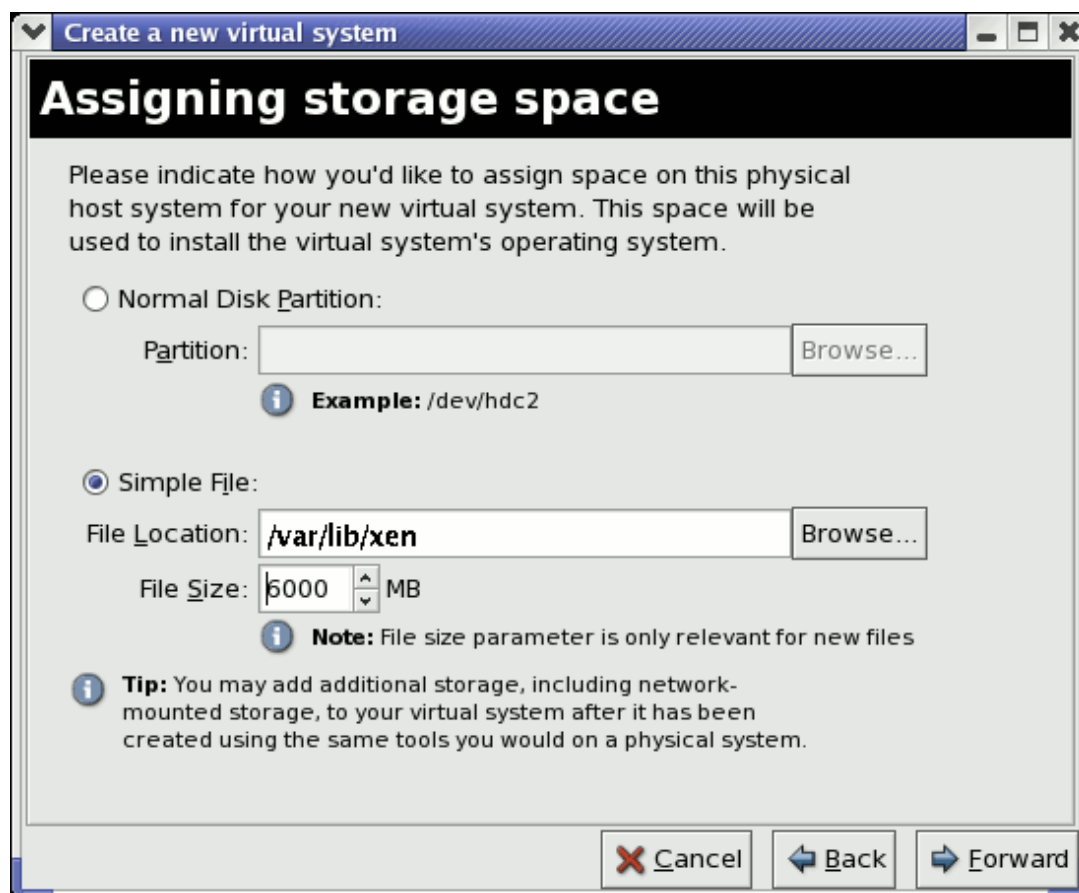


Figure 17.11. Assigning the Storage Space

7. Select memory to allocate the guest and the number of virtual CPUs then click **Forward**.

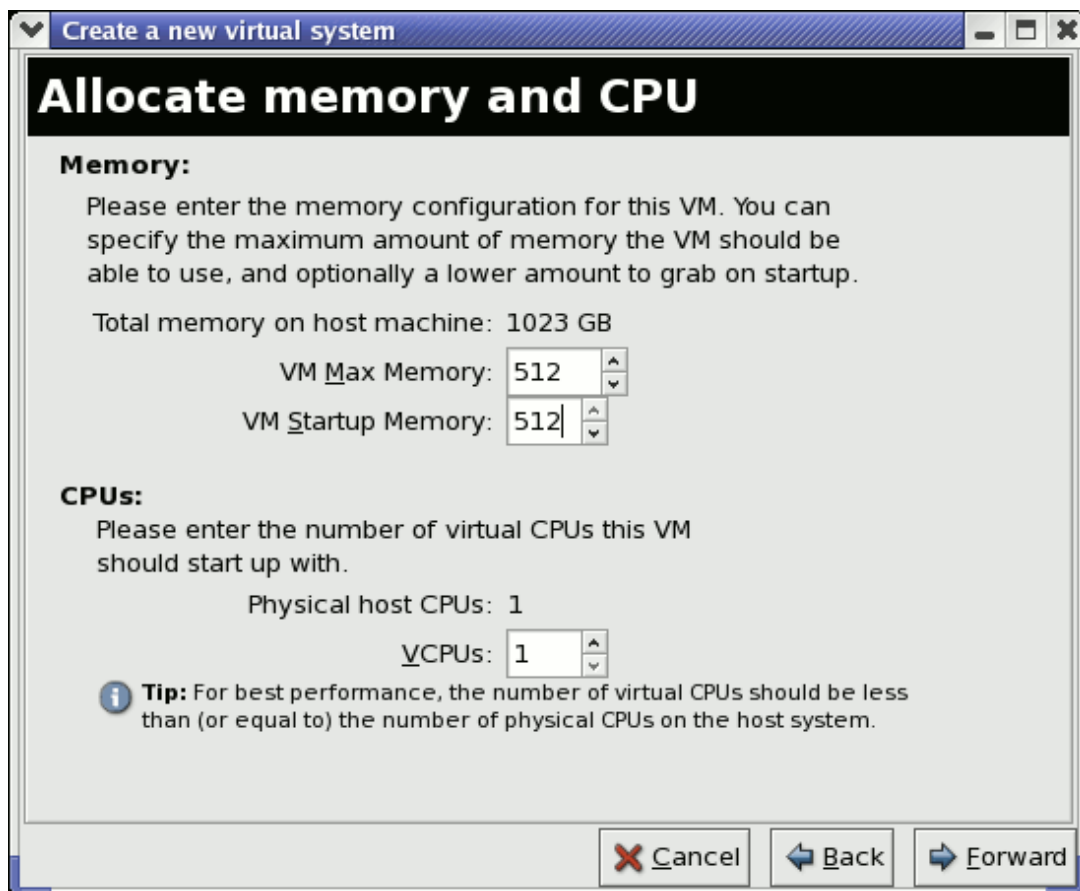


Figure 17.12. Allocating Memory and CPU

8. Select **Forward** to open a console and the files start to install.




Figure 17.13. Allocating Memory and CPU

9. Complete your installation in the window provided.

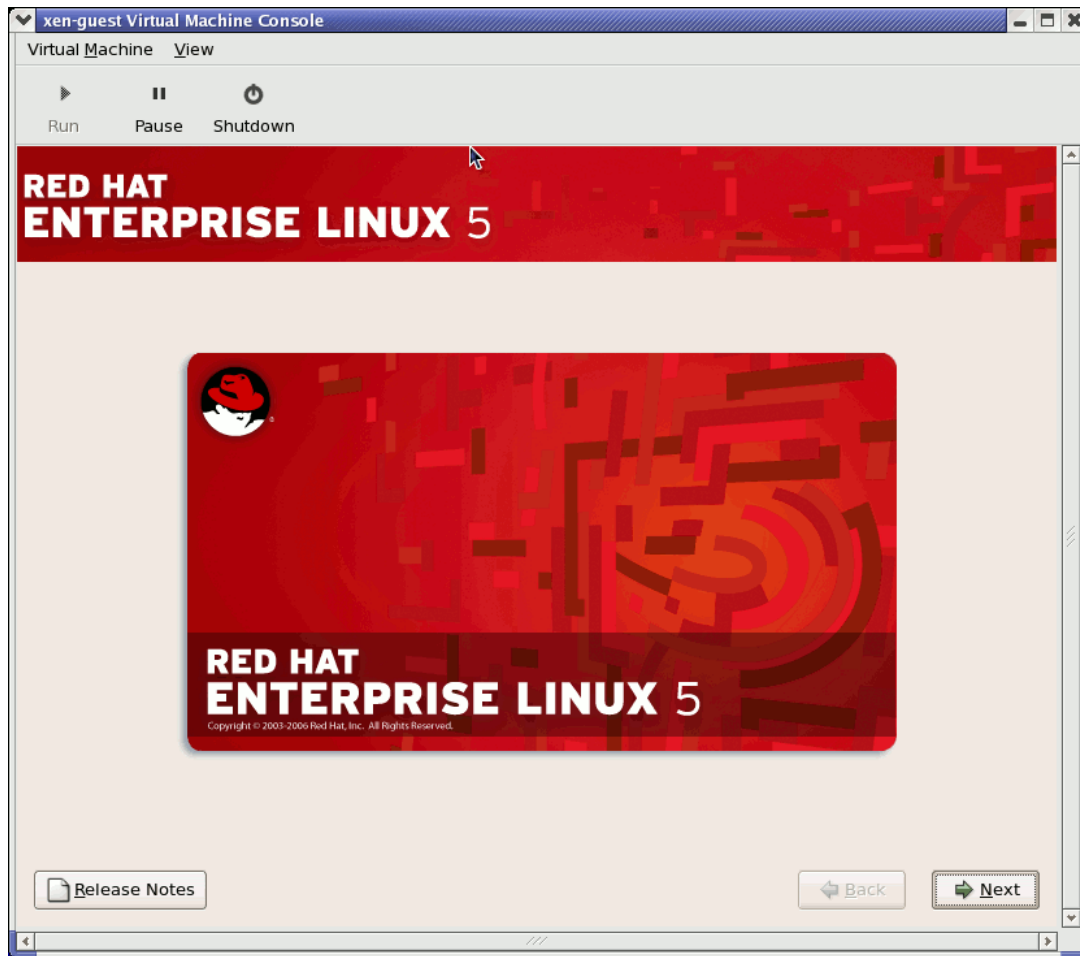


Figure 17.14. Installation Begins...



Warning

When installing Red Hat Enterprise Linux 5.1 on a fully virtualized guest, do *not* use the `kernel-xen` kernel. Using this kernel on fully virtualized guests can cause your system to hang.

If you are using an *Installation Number* when installing Red Hat Enterprise Linux 5.1 on a fully virtualized guest, be sure to deselect the `Virtualization` package group during the installation. The `Virtualization` package group option installs the `kernel-xen` kernel.

Note that paravirtualized guests are not affected by this issue. Paravirtualized guests always use the `kernel-xen` kernel.

10. Type `xm create -c xen-guest` to start the Red Hat Enterprise Linux 5.1 guest. Right click

on the guest in the Virtual Machine Manager and choose **Open** to open a virtual console.




Figure 17.15. Red Hat Enterprise Linux 5.1 (guest)

11. Enter user name and password to continue using the Virtual Machine Manager.

8. Restoring A Saved Machine

After you start the Virtual Machine Manager, all virtual machines on your system are displayed in the main window. Domain0 is your host system. If there are no machines present, this means that currently there are no machines running on the system.

To restore a previously saved session:

1. From the **File** menu, select **Restore a saved machine**.

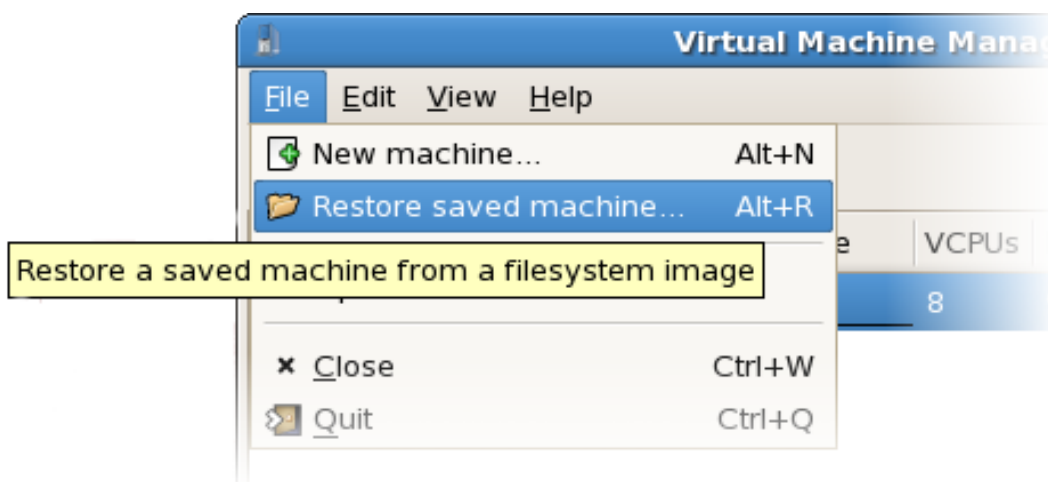


Figure 17.16. Restoring a Virtual Machine

2. The Restore Virtual Machine main window appears.

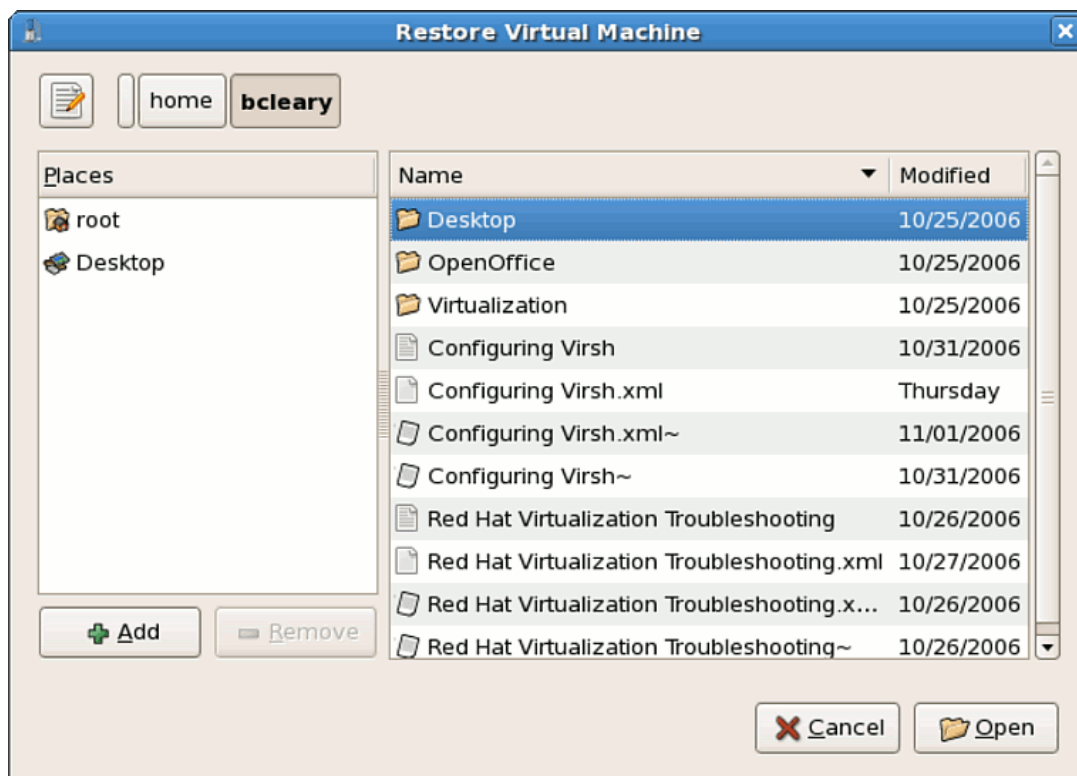


Figure 17.17. Selecting Saved Virtual Machine Session

3. Navigate to correct directory and select the saved session file.
4. Click **Open**.

The saved virtual system appears in the Virtual Machine Manager main window.




Figure 17.18. The Restored Virtual Machine Manager Session

9. Displaying Virtual Machine Details

You can use the Virtual Machine Monitor to view activity data information for any virtual machines on your system.

To view a virtual system's details:

1. In the Virtual Machine Manager main window, highlight the virtual machine that you want to view.




Figure 17.19. Selecting Virtual Machine to Display

- From the Virtual Machine Manager **Edit** menu, select **Machine Details** (or click the **Details** button on the bottom of the Virtual Machine Manager main window).




Figure 17.20. Displaying Virtual Machine Details Menu

The Virtual Machine Details Overview window appears. This window summarizes CPU and memory usage for the domain(s) you specified.

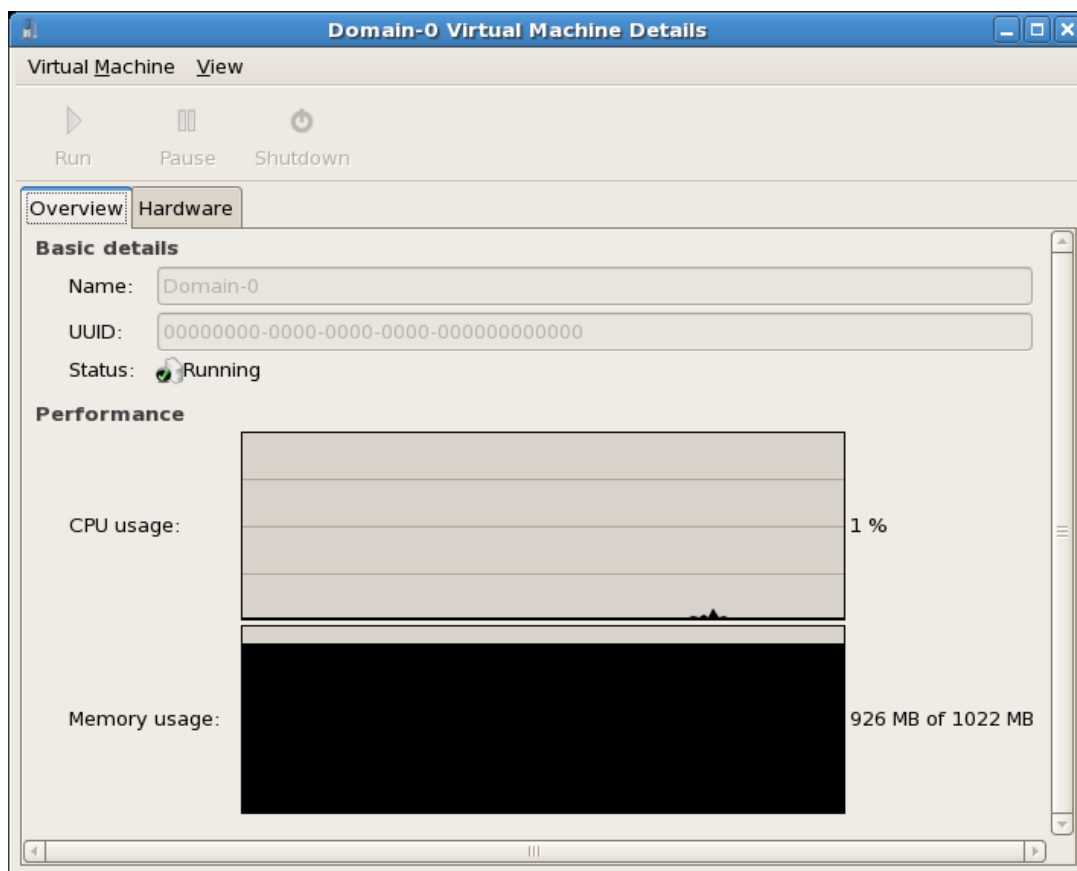


Figure 17.21. Displaying Virtual Machine Details Overview

3. On the Virtual Machine Details window, click the **Hardware** tab.

The Virtual Machine Details Hardware window appears.

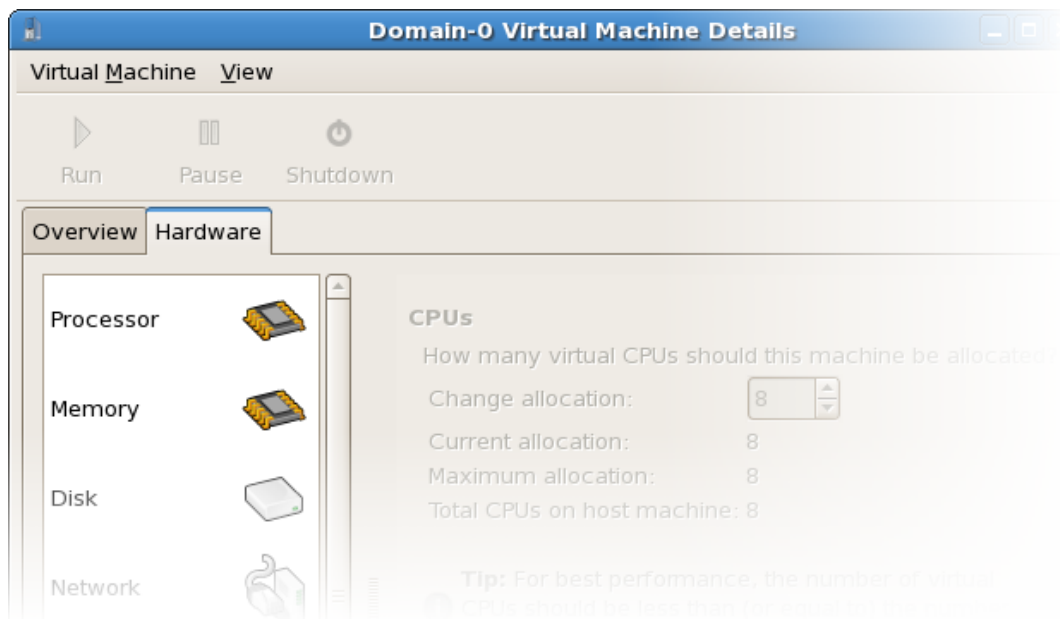


Figure 17.22. Displaying Virtual Machine Details Hardware

4. On the **Hardware** tab, click on **Processor** to view or change the current processor memory allocation.

Figure 17.23. Displaying Processor Allocation

5. On the **Hardware** tab, click on **Memory** to view or change the current RAM memory allocation.

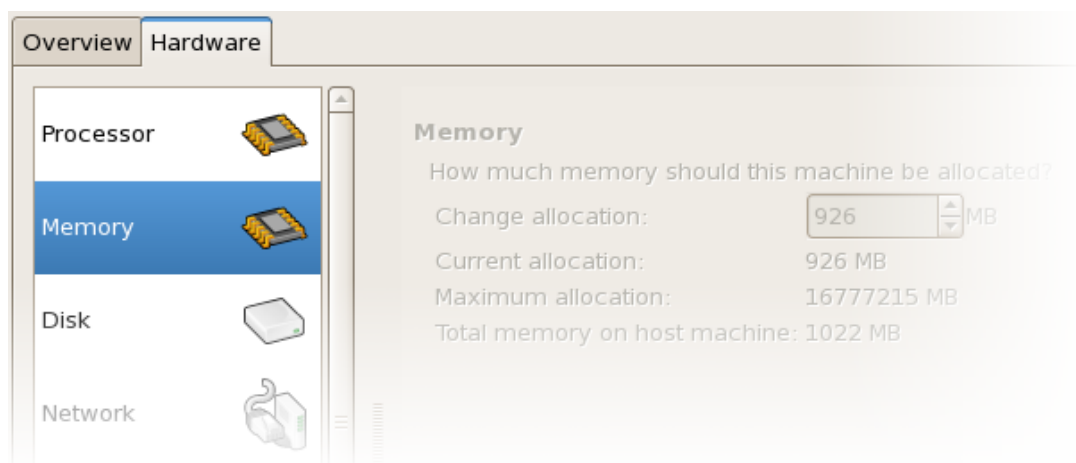


Figure 17.24. Displaying Memory Allocation

6. On the **Hardware** tab, click on **Disk** to view or change the current hard disk configuration.

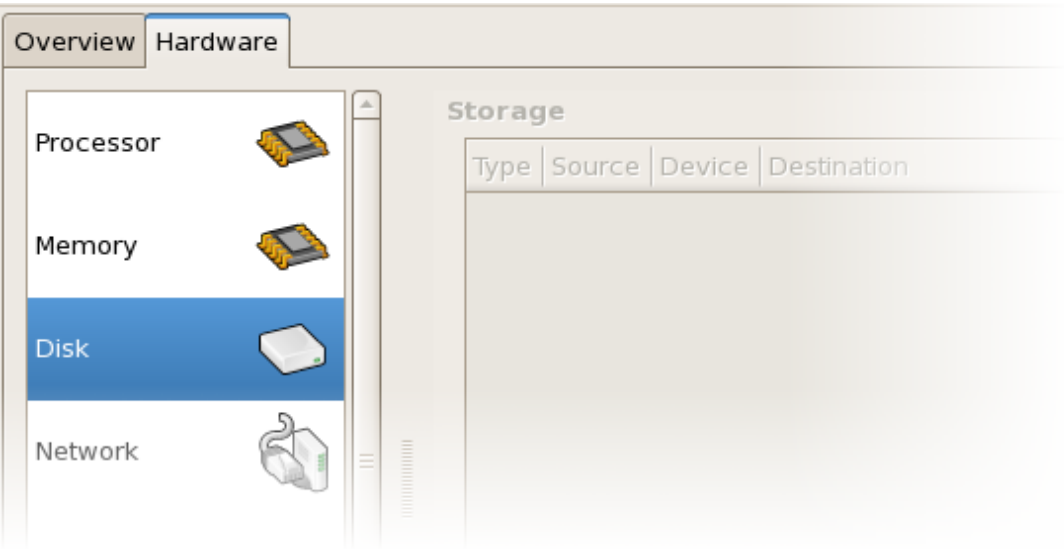


Figure 17.25. Displaying Disk Configuration

7. On the **Hardware** tab, click on **Network** to view or change the current network configuration.

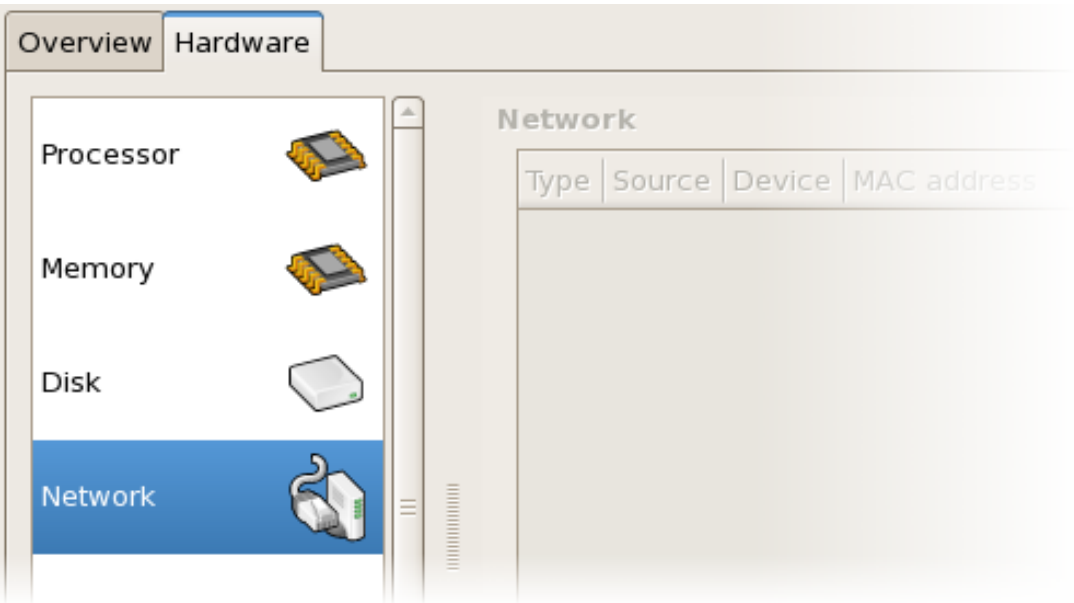



Figure 17.26. Displaying Network Configuration

10. Configuring Status Monitoring

You can use the Virtual Machine Manager to modify the virtual system Status monitoring.

To configure Status monitoring, and enable Consoles:

1. From the **Edit** menu, select **Preferences**.

**Figure 17.27. Modifying Virtual Machine Preferences**

The Virtual Machine Manager Preferences window appears.

2. From the Status monitoring area selection box, specify the time (in seconds) that you want the system to update.

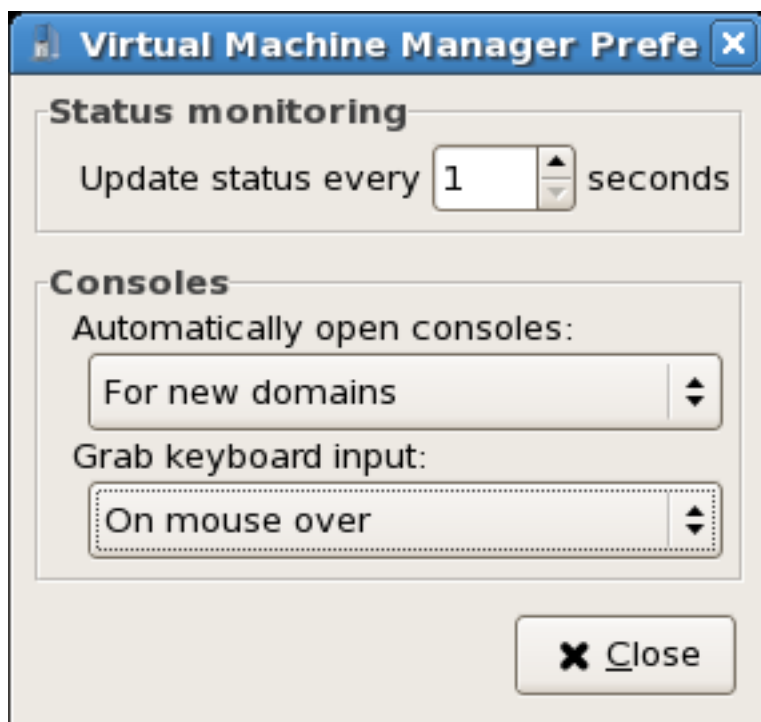


Figure 17.28. Configuring Status Monitoring

3. From the Consoles area, specify how to open a console and specify an input device.

11. Displaying Domain ID

To view the domain IDs for all virtual machines on your system:

1. From the **View** menu, select the **Domain ID** check box.

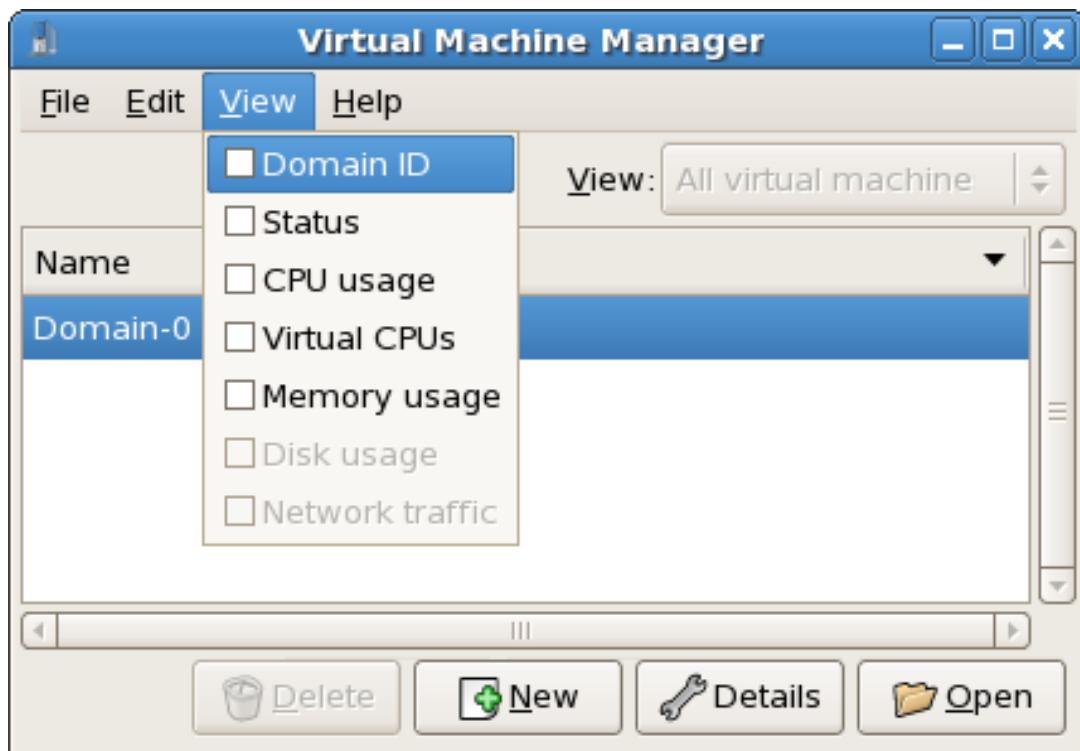


Figure 17.29. Displaying Domain-IDs

2. The Virtual Machine Manager lists the Domain ID's for all domains on your system.




Figure 17.30. Displaying Domain-IDs

12. Displaying Virtual Machine Status

To view the status of all virtual machines on your system:

1. From the **View** menu, select the **Status** check box.

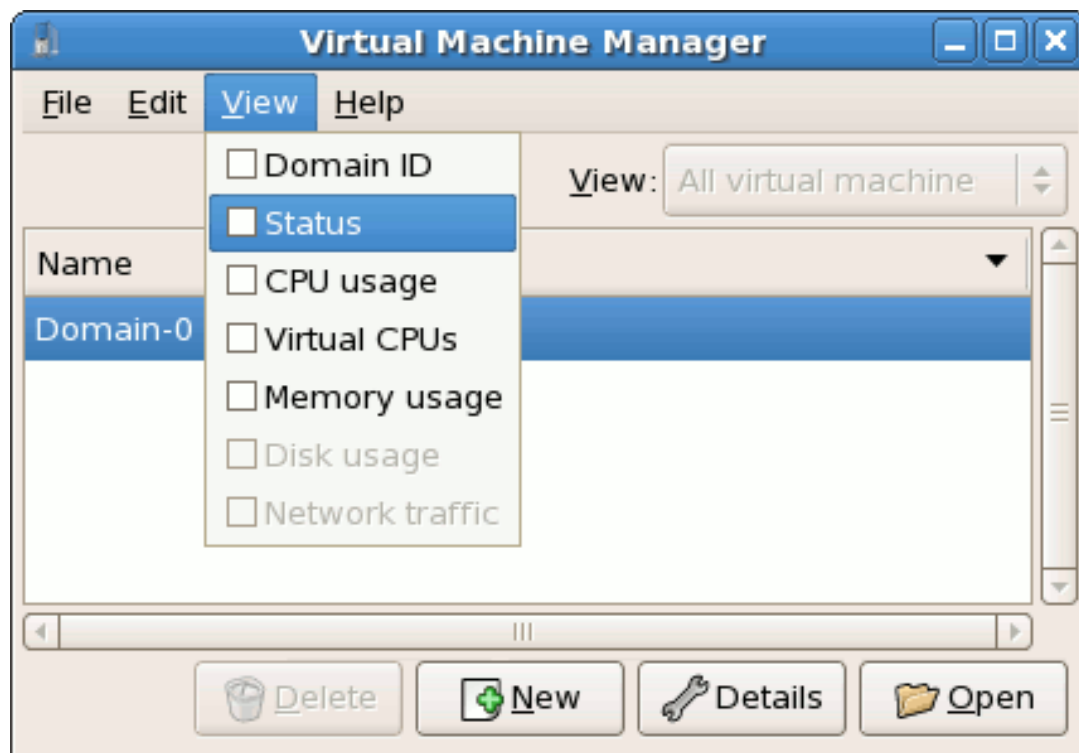


Figure 17.31. Displaying Virtual Machine Status

2. The Virtual Machine Manager lists the status of all virtual machines on your system.




Figure 17.32. Displaying Virtual Machine Status

13. Displaying Virtual CPUs

To view the amount of virtual CPUs for all virtual machines on your system:

1. From the **View** menu, select the **Virtual CPUs** check box.

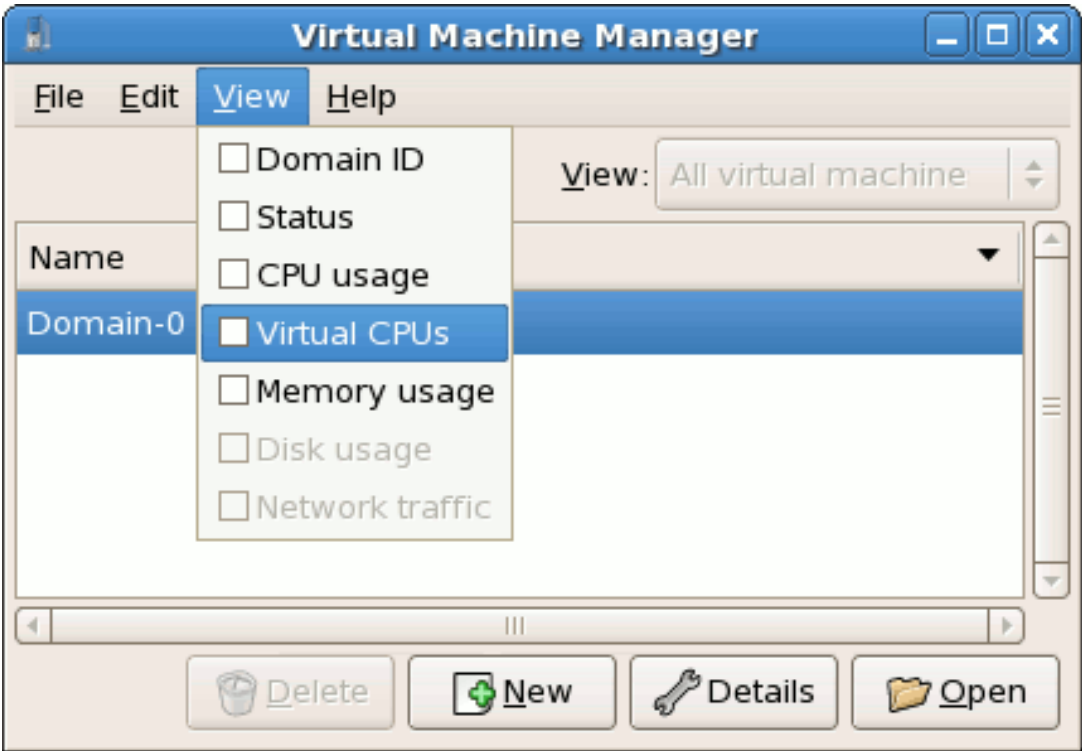


Figure 17.33. Displaying Virtual CPUs

2. The Virtual Machine Manager lists the Virtual CPUs for all virtual machines on your system.


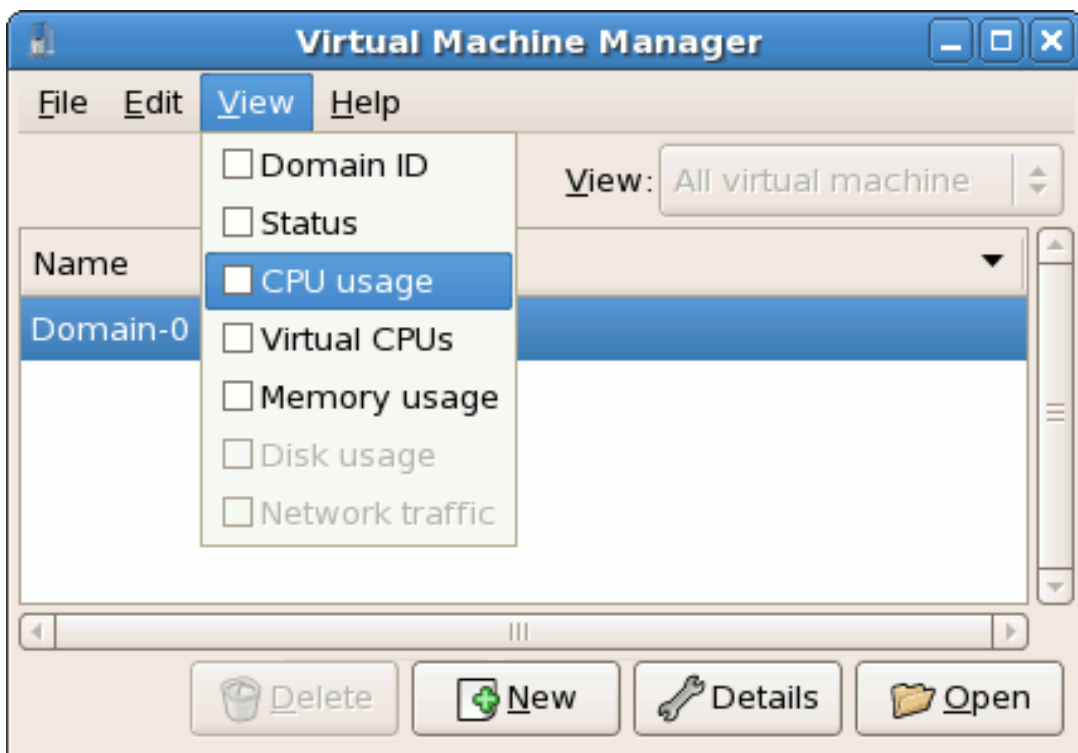


Figure 17.34. Displaying Virtual CPUs

14. Displaying CPU Usage

To view the CPU usage for all virtual machines on your system:

1. From the **View** menu, select the **CPU Usage** check box.

**Figure 17.35. Displaying CPU Usage**

2. The Virtual Machine Manager lists the percentage of CPU in use for all virtual machines on your system.




Figure 17.36. Displaying CPU Usage

15. Displaying Memory Usage

To view the memory usage for all virtual machines on your system:

1. From the **View** menu, select the **Memory Usage** check box.

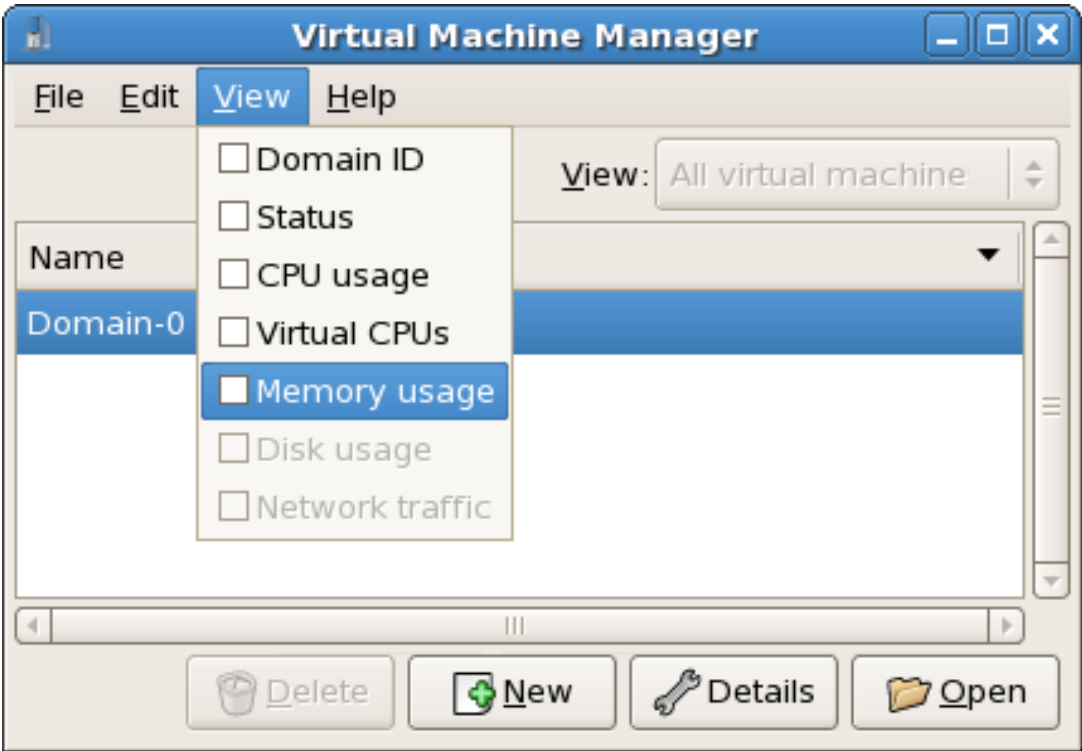


Figure 17.37. Displaying Memory Usage

2. The Virtual Machine Manager lists the percentage of memory in use (in megabytes) for all virtual machines on your system.




Figure 17.38. Displaying Memory Usage

16. Managing a Virtual Network

To configure a virtual network on your system:

1. From the **Edit** menu, select **Host Details**.




Figure 17.39. Selecting Host Details

2. This will open the **Host Details** menu. Click the **Virtual Networks** tab.

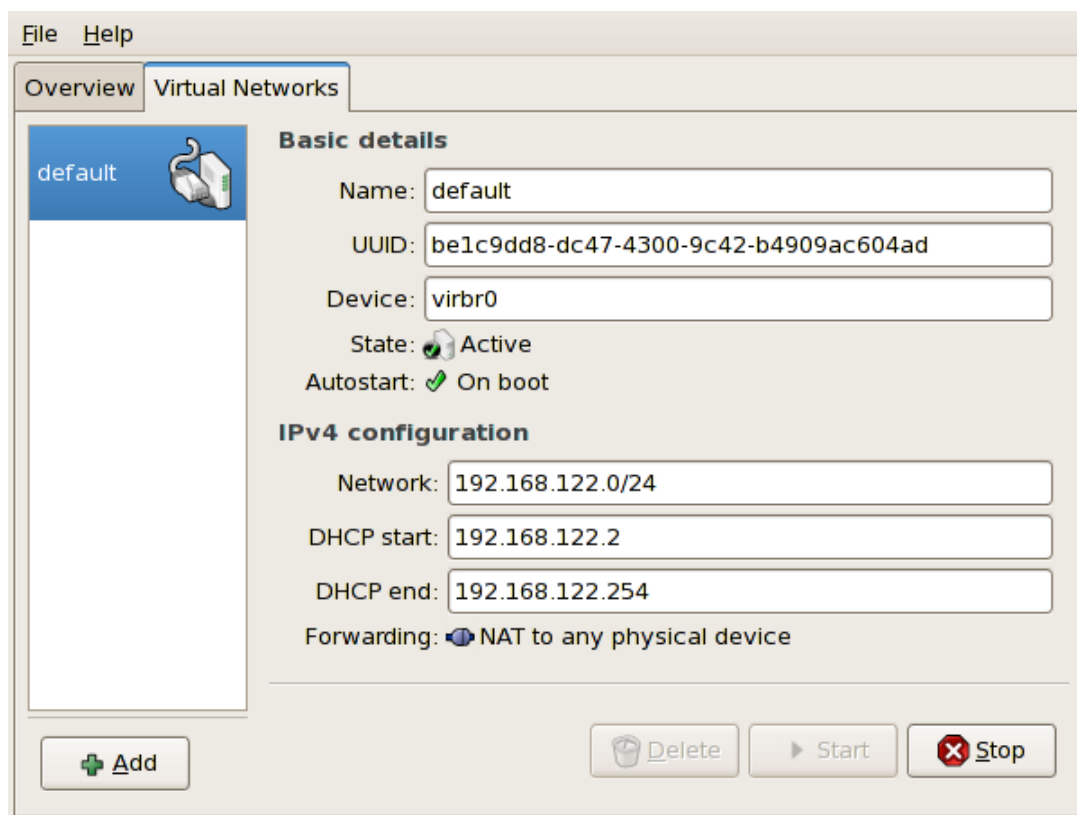


Figure 17.40. Virtual Network Configuration

3. All available virtual networks are listed on the left-hand box of the menu. You can edit the configuration of a virtual network by selecting it from this box and editing as you see fit.

17. Creating a Virtual Network

To create a virtual network on your system:

1. Open the **Host Details** menu (refer to [Section 16, "Managing a Virtual Network"](#)) and click the **Add** button.




Figure 17.41. Virtual Network Configuration

This will open the **Create a new virtual network** menu. Click **Forward** to continue.

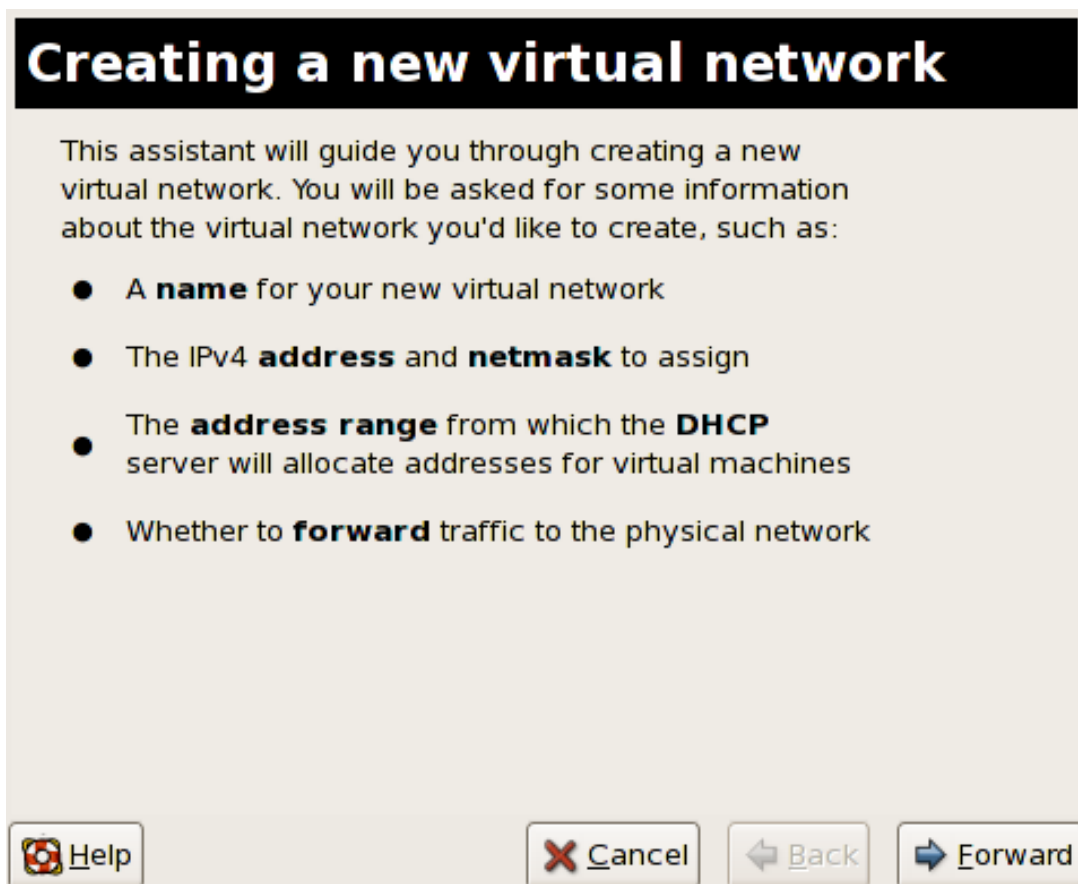


Figure 17.42. Creating a new virtual network

2. Enter an appropriate name for your virtual network and click **Forward**.

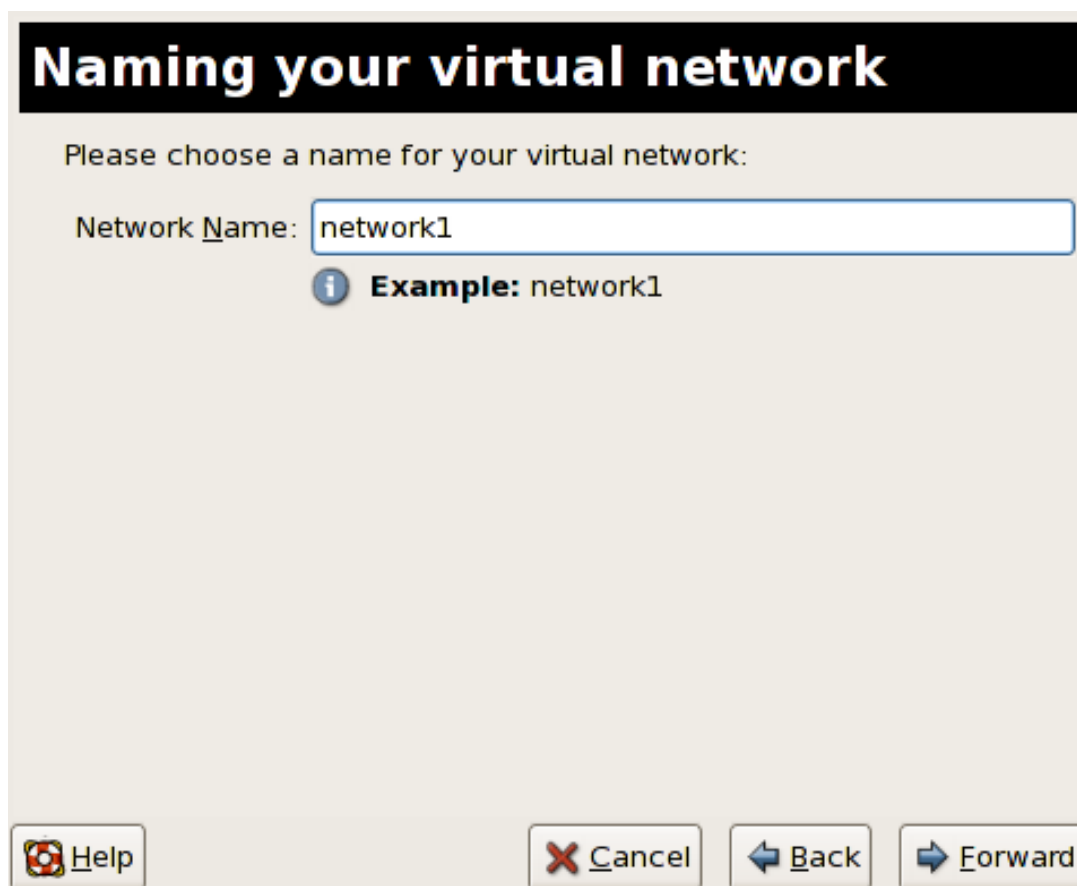



Figure 17.43. Naming your virtual network

3. Enter an IPv4 address space for your virtual network and click **Forward**.

Choosing an IPv4 address space

You will need to choose an IPv4 address space for the virtual network:

Network:

 **Hint:** The network should be chosen from one of the IPv4 private address ranges. eg 10.0.0.0/8, 172.16.0.0/12, or 192.168.0.0/16

Netmask: 255.255.255.0
Broadcast: 192.168.100.255
Gateway: 192.168.100.1
Size: 256 addresses
Type: Private





 Help  Cancel  Back  Forward

Figure 17.44. Choosing an IPv4 address space


4. Define the DHCP range for your virtual network by specifying a **Start** and **End** range of IP addresses. Click **Forward** to continue.

Selecting the DHCP range

Please choose the range of addresses the DHCP server can use to allocate to guests attached to the virtual network

Start:

End:

 **Tip:** Unless you wish to reserve some addresses to allow static network configuration in virtual machines, these parameters can be left with their default values.





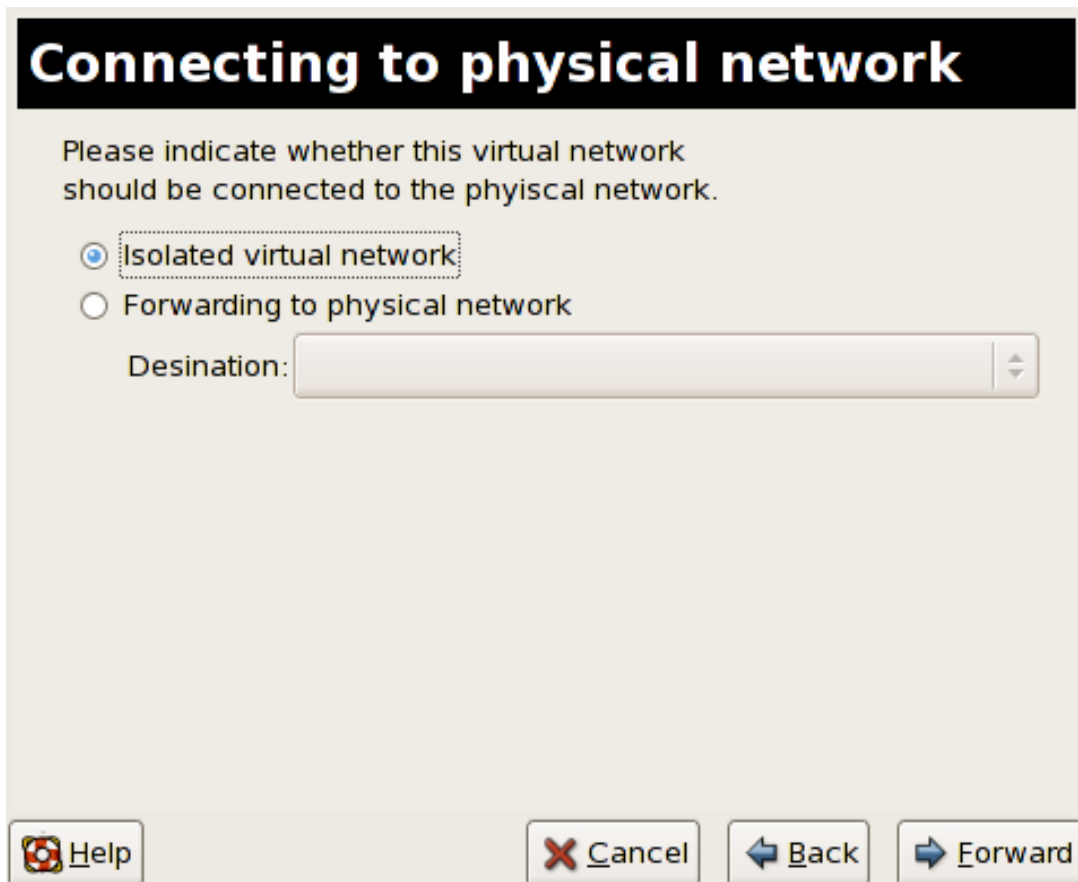
 Help  Cancel  Back  Forward

Figure 17.45. Selecting the DHCP range

5. Select how the virtual network should connect to the physical network.



Connecting to physical network

Please indicate whether this virtual network should be connected to the physical network.

☒ Isolated virtual network

☐ Forwarding to physical network

Desination:





 Help  Cancel  Back  Forward

Figure 17.46. Connecting to physical network

If you select **Forwarding to physical network**, choose whether the **Destination** should be **NAT to any physical device** or **NAT to physical device eth0**.

Click **Forward** to continue.

6. You are now ready to create the network. Check the configuration of your network and click **Finish**.

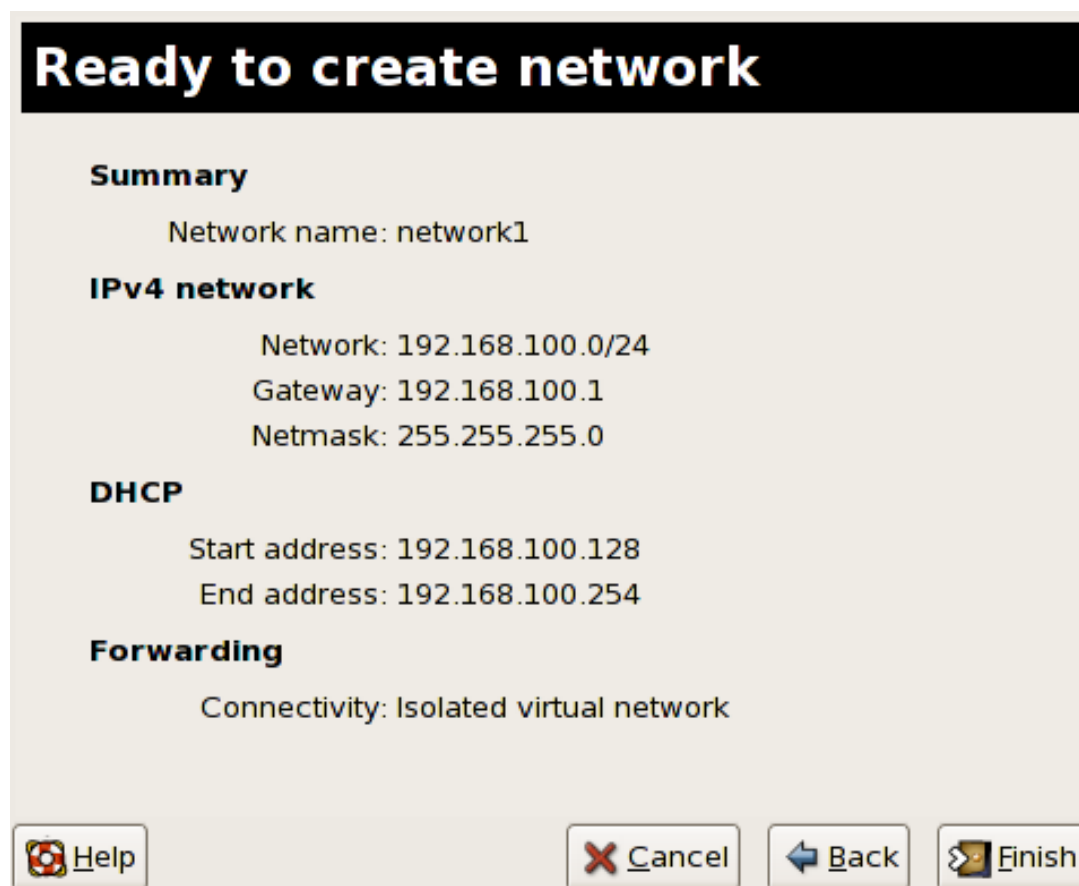


Figure 17.47. Ready to create network

7. The new virtual network is now available in the **Virtual Network** tab of the **Host Details** menu.

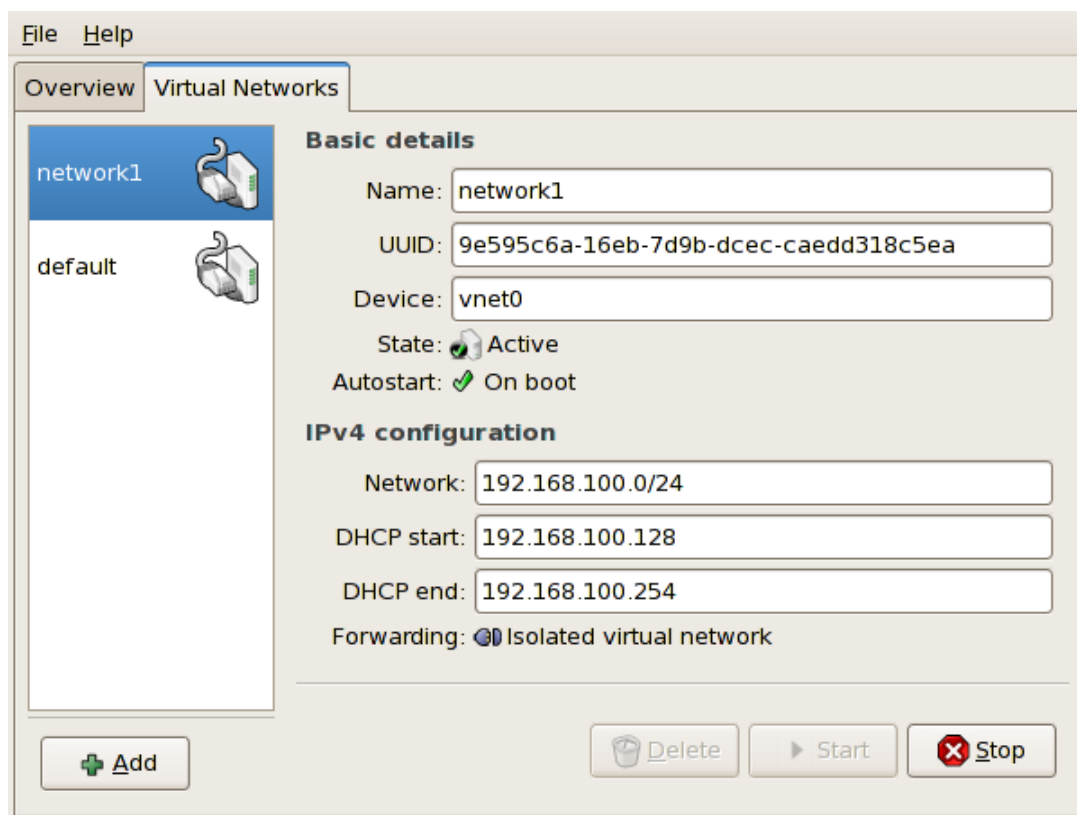


Figure 17.48. New virtual network is now available

Red Hat Virtualization

Troubleshooting

This section covers potential issues you may experience in the installation, management, and general day-to-day operations of your Red Hat Virtualization system(s). This troubleshooting section covers the error messages, log file locations, system tools, and general approaches to research data and analyze problems.

1. Logfile Overview and Locations

When deploying Red Hat Enterprise Linux 5.1 with Virtualization into your network infrastructure, the host's Virtualization software uses many specific directories for important configuration, log files, and other utilities. All the Red Hat Virtualization logs files are standard ASCII files, and easily accessible with any ASCII based editor:

- The Red Hat Virtualization main configuration directory is `/etc/xen/`. This directory contains the `xend` daemon and other virtual machine configuration files. The networking script files reside here as well (in the `/scripts` subdirectory).
- All of actual log files themselves that you will consult for troubleshooting purposes reside in the `/var/log/xen` directory.
- You should also know that the default directory for all virtual machine file-based disk images resides in the `/var/lib/xen` directory.
- Red Hat Virtualization information for the `/proc` file system reside in the `/proc/xen/` directory.

2. Logfile Descriptions

Red Hat Virtualization features the `xend` daemon and `qemu-dm` process, two utilities that write the multiple log files to the `/var/log/xen/` directory:

- `xend.log` is the logfile that contains all the data collected by the `xend` daemon, whether it is a normal system event, or an operator initiated action. All virtual machine operations (such as create, shutdown, destroy, etc.) appears here. The `xend.log` is usually the first place to look when you track down event or performance problems. It contains detailed entries and conditions of the error messages.
- `xend-debug.log` is the logfile that contains records of event errors from `xend` and the Virtualization subsystems (such as framebuffer, Python scripts, etc.).

- `xen-hotplug-log` is the logfile that contains data from hotplug events. If a device or a network script does not come online, the event appears here.
- `qemu-dm.[PID].log` is the logfile created by the `qemu-dm` process for each fully virtualized guest. When using this logfile, you must retrieve the given `qemu-dm` process PID, by using the `ps` command to examine process arguments to isolate the `qemu-dm` process on the virtual machine. Note that you must replace the `[PID]` symbol with the actual PID `qemu-dm` process.

If you encounter any errors with the Virtual Machine Manager, you can review the generated data in the `virt-manager.log` file that resides in the `/var/lib/libvirt` directory. Note that every time you start the Virtual Machine Manager, it overwrites the existing logfile contents. Make sure to backup the `virt-manager.log` file, before you restart the Virtual Machine manager after a system error.

3. Important Directory Locations

There are additional utilities and logfiles you should remember when you track errors and troubleshoot problems within Red Hat Virtualization environments:

- Virtual machines images reside in the `/var/lib/xen/images` directory.
- When you restart the `xend` daemon, it updates the `xend-database` that resides in the `/var/lib/xen/xend-db` directory.
- Virtual machine dumps (that you perform with `xm dump-core` command) resides in the `/var/lib/xen/dumps` directory.
- The `/etc/xen` directory contains the configuration files that you use to manage system resources. The `xend` daemon configuration file is called `xend-config.sxp` and you can use this file to implement system-wide changes and configure the networking callouts.
- The `proc` commands are another resource that allows you to gather system information. These `proc` entries reside in the `/proc/xen` directory:

`/proc/xen/capabilities`

`/proc/xen/balloon`

`/proc/xen/xenbus/`

4. Troubleshooting Tools

This section summarizes the System Administrator applications, the networking utilities, and the Advanced Debugging Tools (for more information on using these tools to configure the Red Hat Virtualization services, see the respective configuration documentation). You can employ these standard System Administrator Tools and logs to assist with troubleshooting:

- `xentop`
- `xm dmesg`
- `xm log`
- `vmstat`
- `iostat`
- `lsof`

You can employ these Advanced Debugging Tools and logs to assist with troubleshooting:

- `XenOprofile`
- `systemTap`
- `crash`
- `sysrq`
- `sysrq t`
- `sysrq w`

You can employ these Networking Tools to assist with troubleshooting:

- `ifconfig`
- `tcpdump`
- `brctl`

`brctl` is a networking tool that inspects and configures the ethernet bridge configuration in the Virtualization linux kernel. You must have root access before performing these example commands:

```
# brctl show

bridge-name    bridge-id      STP   enabled  interfaces
-----
```

```
xenbr0          8000.feffffff    no      vif13.0
xenbr1          8000.ffffefff    yes     pddummy0
xenbr2          8000.fffffefe    no      vif0.0

# brctl showmacs xenbr0

port-no         mac-addr                local?    ageing timer

1               fe:ff:ff:ff:ff:         yes       0.00
2               fe:ff:ff:fe:ff:         yes       0.00

# brctl showstp xenbr0

xenbr0

bridge-id       8000.feffffffff

designated-root  8000.feffffffff

root-port       0                      path-cost  0

max-age         20.00                 bridge-max-age  20.00

hello-time      2.00                  bridge-hello-time  2.00

forward-delay   0.00                  bridge-forward-delay  0.00

ageing-time     300.01

hello-timer     1.43                  tcu-timer    0.00

topology-change-timer 0.00                gc-timer     0.02
```

5. Troubleshooting with the Logs

When encountering issues with installing Red Hat Virtualization, you can refer to the host system's two logs to assist with troubleshooting. The `xend.log` file contains the same basic information as when you run the `xm log` command. It resides in the `/var/log/` directory. Here is an example log entry for when you create a domain running a kernel:

```
[2006-12-27 02:23:02 xend] ERROR (SrvBase: 163) op=create: Error creating
domain: (0, 'Error')
Traceback (most recent call list)
File "/usr/lib/python2.4/site-packages/xen/xend/server/SrvBase.py" line 107
in_perform val = op_method (op,req)
File
"/usr/lib/python2.4/site-packages/xen/xend/server/SrvDomainDir.py line 71 in
op_create
raise XendError ("Error creating domain: " + str(ex))
XendError: Error creating domain: (0, 'Error')
```

The other log file, `xend-debug.log`, is very useful to system administrators since it contains even more detailed information than `xend.log`. Here is the same error data for the same kernel domain creation problem:

```
ERROR: Will only load images built for Xen v3.0
ERROR: Actually saw: GUEST_OS=netbsd, GUEST_VER=2.0, XEN_VER=2.0;
LOADER=generic, BSD_SYMTAB'
ERROR: Error constructing guest OS
```

When calling customer support, always include a copy of both these log files when contacting the technical support staff.

6. Troubleshooting with the Serial Console

The serial console is helpful in troubleshooting difficult problems. If the Virtualization kernel crashes and the hypervisor generates an error, there is no way to track the error on a local host. However, the serial console allows you to capture it on a remote host. You must configure the Xen host to output data to the serial console. Then you must configure the remote host to capture the data. To do this, you must modify these options in the `grub.conf` file to enable a 38400-bps serial console on `com1 /dev/ttyS0`:

```
title Red Hat Enterprise Linux (2.6.18-8.2080_RHEL5xen0)
    root (hd0,2)
    kernel /xen.gz-2.6.18-8.el5 com1=38400,8n1
    module /vmlinuz-2.6.18-8.el5xen ro root=LABEL=/rhgb quiet
console=xvc console=tty xencons=xvc
    module /initrd-2.6.18-8.el5xen.img
```

The `sync_console` can help determine a problem that causes hangs with asynchronous hypervisor console output, and the `"pnpcapi=off"` works around a problem that breaks input on the serial console. The parameters `"console=ttyS0"` and `"console=tty"` means that kernel errors get logged with on both the normal VGA console and on the serial console. Then you can install and set up `ttymwatch` to capture the data on a remote host connected by a standard null-modem cable. For example, on the remote host you could type:

```
ttymwatch --name myhost --port /dev/ttyS0
```

This pipes the output from `/dev/ttyS0` into the file `/var/log/ttymwatch/myhost.log`.

7. Paravirtualized Guest Console Access

Paravirtualized guest operating systems automatically has a virtual text console configured to plumb data to the Domain0 operating system. You can do this from the command line by typing:

```
xm console [domain name or number]
```

Where `domain100` represents a running name or number. You can also use the Virtual Machine Manager to display the virtual text console. On the Virtual Machine Details window, select **Serial Console** from the **View** menu.

8. Full Virtualization Guest Console Access

Full Virtualized guest operating systems automatically has a text console configured for use, but the difference is the kernel guest is not configured. To enable the guest virtual serial console to work with the Full Virtualized guest, you must modify the guest's `grub.conf` file, and include the `'console =ttyS0 console=tty0'` parameter. This ensures that the kernel messages are sent to the virtual serial console (and the normal graphical console). If you plan to use the virtual serial console in a full virtualized guest, you must edit the configuration file in the `/etc/xen/` directory. On the host domain, you can then access the text console by typing:

```
xm console
```

You can also use the Virtual Machine Manager to display the serial console. On the Virtual Machine Details window, select Serial Console from the View menu.

9. Implementing Lun Persistence

If your system is not using multipath, you can use `udev` to implement lun persistence. Before implementing lun persistence in your system, ensure that you acquire the proper UUIDs. Once you acquire these, you can configure lun persistence by editing the `scsi_id` file that resides in the `/etc` directory. Once you have this file open in a text editor, you must comment out this line:

```
# options=-b
```

Then replace it with this parameter:

```
# options=-g
```

This tells `udev` to monitor all system SCSI devices for returning UUIDs. To determine the system UUIDs, type:

```
# scsi_id -g -s /block/sdc
```

The output should resemble the following:

```
[root@devices] # scsi_id -g -s /block/sdc
*3600a0b80001327510000015427b625e*
```

This long string of characters is the UUID. To get the device names to key off the UUID, check each device path to ensure that the UUID number is the same for each device. The UUIDs do not change when you add a new device to your system. Once you have checked the device paths, you must create rules for the device naming. To create these rules, you must edit the `20-names.rules` file that resides in the `/etc/udev/rules.d` directory. The device naming rules you create here should follow this format:

```
# KERNEL="sd*", BUS="scsi", PROGRAM="sbin/scsi_id", RESULT="UUID",
NAME="devicename"
```

Replace your existing UUID and devicename with the above UUID retrieved entry. So the rule should resemble the following:

```
KERNEL="sd*", BUS="scsi", PROGRAM="sbin/scsi_id",
RESULT="3600a0b80001327510000015427b625e
", NAME="mydevicename"
```

This causes the system to enable all devices that match `/dev/sd*` to inspect the given UUID. When it finds a matching device, it creates a device node called `/dev/devicename`. For this example, the device node is `/dev/mydevice`. Finally, you need to append the `rc.local` file that resides in the `/etc` directory with this path:

```
/sbin/start_udev
```

IMPLEMENTING LUN PERSISTENCE WITH MULTIPATH

To implement lun persistence in a multipath environment, you must define the alias names for the multipath devices. For this example, you must define four device aliases by editing the `multipath.conf` file that resides in the `/etc/` directory:

```
multipath {
    wwid      3600a0b80001327510000015427b625e
    alias     oraml
}
multipath {
    wwid      3600a0b80001327510000015427b6
```

```
        alias      oramp2
    }
    multipath {
        wwid      3600a0b80001327510000015427b625e
        alias      oramp3
    }
    multipath {
        wwid      3600a0b80001327510000015427b625e
        alias      oramp4
    }
}
```

This defines 4 luns: `/dev/mpath/oramp1`, `/dev/mpath/oramp2`, `/dev/mpath/oramp3`, and `/dev/mpath/oramp4`. The devices will reside in the `/dev/mpath` directory. These lun names are persistent over reboots as it creates the alias names on the wwid of the luns.

10. SELinux Considerations

This sections contains things to you must consider when you implement SELinux into your Red Hat Virtualization environment. When you deploy system changes or add devices, you must update your SELinux policy accordingly. To configure an LVM volume for a guest, you must modify the SELinux context for the respective underlying block device and volume group.

```
# semanage fcontext -a -t xen_image _t -f -b /dev/sda2
# restorecon /dev/sda2
```

The boolean parameter `xend_disable_trans` put `xend` in unconfined mode after restarting the daemon. It is better to disable protection for a single daemon than the whole system. It is advisable that you should not re-label directories as `xen_image_t` that you will use elsewhere.

11. Accessing Data on Guest Disk Image

You can use two separate applications that assist you in accessing data from within a guest disk image. Before using these tools, you must shut down the guests. Accessing the file system from the guest and dom0 could potentially harm your system.

You can use the `kpartx` application to handle partitioned disks or LVM volume groups:

```
yum install kpartx
kpartx -av /dev/xen/guest1
add map guest1p1 : 0 208782 linear /dev/xen/guest1 63
add map guest1p2: 0 16563015 linear /dev/xen/guest1 208845
```

To access LVM volumes on a second partition, you must rescan LVM with `vgscan` and activate the volume group on the partition (called `VolGroup00` by default) by using the `vgchange -ay` command:

```
# kpartx -a /dev/xen/guest1
#vgscan
Reading all physical volumes . This may take a while...
Found volume group "VolGroup00" using metadata type lvm2
# vgchange -ay VolGroup00
2 logical volume(s) in volume group VolGroup00 now active.
# lvs
LV VG Attr Lsize Origin Snap% Move Log Copy%
LogVol00 VolGroup00 -wi-a- 5.06G
LogVol01 VolGroup00 -wi-a- 800.00M
# mount /dev/VolGroup00/LogVol00 /mnt/
....
#umount /mnt/
#vgchange -an VolGroup00
#kpartx -d /dev/xen/guest1
```

You must remember to deactivate the logical volumes with `vgchange -an`, remove the partitions with `kpartx -d`, and delete the loop device with `losetup -d` when you finish.

12. Common Troubleshooting Situations

When you attempt to start the `xend` service nothing happens. You type `xm listl` and receive the following:

```
Error: Error connecting to xend: Connection refused. Is xend running?
```

You try to run `xend start` manually and receive more errors:

```
Error: Could not obtain handle on privileged command interfaces (2 = No such
file or directory)
Traceback (most recent call last:)

File "/usr/sbin/xend/", line 33 in ?

from xen.xend.server import SrvDaemon

File "/usr/lib/python2.4/site-packages/xen/xend/server/SrvDaemon.py" , line
26 in ?

from xen.xend import XendDomain

File "/usr//lib/python2.4/site-packages/xen/xend/XendDomain.py" , line 33,
in ?

from xen.xend import XendDomainInfo

File "/usr/lib/python2.4/site-packages/xen/xend/image.py" , line37, in ?

import images
```

```
File "/usr/lib/python2.4/site-packages/xen/xend/image.py" , line30, in ?
xc = xen.lowlevel.xc.xc ()

RuntimeError: (2, 'No such file or directory' )
```

What is most likely happened here is that you rebooted your host into a kernel that is not a `xen-hypervisor` kernel. To correct this, you must select the `xen-hypervisor` kernel at boot time (or set the `xen-hypervisor` kernel to default in your `grub.conf` file).

13. Loop Device Errors

If you use file-based guest images, one may have increased the number of configured loop devices (the default allows up to 8 loop devices to become active). If you need more than 8 file-based guests/loop devices, you must modify the `/etc/modprobe.conf` file. When modifying the `modprobe.conf` file, you must include this line:

```
options loop max_loop=64
```

This example uses 64 but you can specify another number to set the maximum loop value. You may also have to implement loop device backed guests on your system. To employ loop device backed guests for a paravirtual system, use the `phy: block device` or `tap:aio` commands. To employ loop device backed guests for a full virtualized system, use the `phy: device` or `file: file` commands.

14. Guest Creation Errors

When you attempt to create a guest, you receive an "Invalid argument" error message. This usually means that the kernel image you are trying to boot is incompatible with the hypervisor. An example of this would be if you were attempting to run a non-PAE FC5 kernel on a PAE only FC6 hypervisor.

You do a yum update and receive a new kernel, the `grub.conf` default kernel switches right back to a bare-metal kernel instead of the Virtualization kernel.

To correct this problem you must modify the default kernel RPM that resides in the `/etc/sysconfig/kernel/` directory. You must ensure that `kernel-xen` parameter is set as the default option in your `gb.conf` file.

15. Serial Console Errors

You receive no output to the serial console. To correct this problem, you must modify the `grub.conf` and change the com port parameters to:


```
serial --unit=1 --speed=115200
```

```
title RHEL5 i386 Xen (2.6.18-1.2910.el5xen)
root (hd0, 8)
kernel /boot/xen.gz-2.6.18-1.2910.el5 com2=115200, 8n1
module /boot/vmlinuz-2.6.18-1.2910.el5xen to root=LABEL=RHEL5_i386
console=tty console=ttyS1115200
module /boot/initrd-2.8.6.18-12910.el5xen.img

title RHEL5 i386 xen (2.6.18.-1.2910.el5xen
root (hd0, 8)
kernel /boot/xen.gz-2.6.18-1.2910.el5 com2=115200 console=com2l
module /boot/vmlinuz2.6.18-1.2910.el5xen to root=LABEL=RHEL5_i386
console=xvc xencons=xvc
module /boot/ititrd-2.6.18-1.2910.el5xen.img
```

These changes to the `grub.conf` should enable your serial console to work correctly. You should be able to use any number for the `ttyS` and it should work like `ttyS0` .

16. Network Bridge Errors

Red Hat Virtualization can configure multiple Virtualization network bridges to use with multiple ethernet cards. To successfully configure multiple network bridges for ethernet cards, you must configure the second network interface by either using the system-config-network TUI/GUI, or by creating a new configuration file in `/etc/sysconfig/network-scripts` . You should use a process to setup multiple Xen bridges. This is an example config file for a second NIC called 'eth1' :

```
#/etc/sysconfig/network-scripts/fcfg-eth1
DEVICE=eth1
BOOTPROTO=static
ONBOOT=yes
USERCTL=no
IPV6INIT=no
PEERDNS=yes
TYPE=Ethernet
NETMASK=255.255.255.0
IPADDR=10.1.1.1
GATEWAY=10.1.1.254
ARP=yes
```

Copy the `/etc/xen/scripts/network-bridge` to `/etc/xen/scripts/network-bridge.xen` .

Edit `/etc/xen/xend-config.sxp` and add a line to your new network bridge script (this example uses "network-virtualization-multi-bridge").

In the `xend-config.sxp` file, the new line should reflect your new script:

```
network-script network-xen-multi-bridge
```

Make sure to uncomment the line that states:

```
network-script network-bridge
```

If you want to create multiple Xen bridges, you must create a custom script. This example below creates two Xen bridges (called `xenbr0` and `xenbr1`) and attaches them to `eth1` and `eth0` , respectively:

```
# !/bin/sh
# network-xen-multi-bridge
# Exit if anything goes wrong
set -e
# First arg is operation.
OP=$1
shift
script=/etc/xen/scripts/network-bridge.xen
case ${OP} in
start)
$script start vifnum=1 bridge=xenbr1 netdev=eth1
$script start vifnum=0 bridge=xenbr0 netdev=eth0
;;
stop)
$script stop vifnum=1 bridge=xenbr1 netdev=eth1
$script stop vifnum=0 bridge=xenbr0 netdev=eth0
;;
status)
$script status vifnum=1 bridge=xenbr1 netdev=eth1
$script status vifnum=0 bridge=xenbr0 netdev=eth0
;;
*)
echo 'Unknown command: ' ${OP}
echo 'Valid commands are: start, stop, status'
exit 1
esac
```

If you want to create additional bridges, just use the example script and copy/paste the file accordingly.

17. Laptop Configurations

The task of configuring your RHEL 5.1 loaded laptop for use on a network environment, presents a number of potential challenges. Most WiFi and wired connections switch constantly during any given day, and Red Hat Virtualization assumes it has access to the same interface consistently. This results in the system performing ifup/ifdown calls to the network interface in

use by Red Hat Virtualization. WiFi cards are not the ideal network connection method since Red Hat Virtualization uses the default network interface.

The idea here is to create a 'dummy' network interface for Red Hat Virtualization to use.

This technique allows you to use a hidden IP address space for your guests and Virtual Machines. To do this operation successfully, you must use static IP addresses as DHCP does not listen for IP addresses on the dummy network. You also must configure NAT/IP masquerading to enable network access for your guests and Virtual Machines. You should attach a static IP when you create the 'dummy' network interface.

For this example, the interface is called `dummy0` and the IP used is `10.1.1.1`. The script is called `ifcfg-dummy0` and resides in the `/etc/sysconfig/network-scripts/` directory:

```
DEVICE=dummy0
BOOTPROTO=none
ONBOOT=yes
USERCTL=no
IPV6INIT=no
PEERDNS=yes
TYPE=Ethernet
NETMASK=255.255.255.0
IPADDR=10.1.1.1
ARP=yes
```

You should bind `xenbr0` to `dummy0` to allow network connection even when disconnected from the physical network.

You will need to make additional modifications to the `xend-config.sxp` file. You must locate the (`network-script 'network-bridge' bridge=xenbr0`) section and add include this in the end of the line:

```
netdev=dummy0
```

You must also make some modifications to your guest's domU networking configuration to enable the default gateway to point to `dummy0`. You must edit the DomU 'network' file that resides in the `/etc/sysconfig/` directory to reflect the example below:

```
NETWORKING=yes
HOSTNAME=localhost.localdomain
GATEWAY=10.1.1.1

IPADDR=10.1.1.10
NETMASK=255.255.255.0
```

It is a good idea to enable NAT in domain0 so that domU can access the public net. This way, even wireless users can work around the Red Hat Virtualization wireless limitations. To do this, you must modify the `S99XenLaptopNAT` file that resides in the `/etc/rc3.d` directory to reflect the example below:

```
#!/bin/bash
#
# XenLaptopNAT  Startup script for Xen on Laptops
#
# chkconfig: - 99 01
# description: Start NAT for Xen Laptops
#
# PATH=/usr/bin:/sbin:/bin:/usr/sbin
# export PATH
GATEWAYDEV=`ip route | grep default | awk {'print $5'}`
iptables -F
case "$1" in
start)
if test -z "$GATEWAYDEV"; then
echo "No gateway device found"
else
echo "Masquerading using $GATEWAYDEV"
/sbin/iptables -t nat -A POSTROUTING -o $GATEWAYDEV -j MASQUERADE
fi
echo "Enabling IP forwarding"
echo 1 > /proc/sys/net/ipv4/ip_forward
echo "IP forwarding set to `cat /proc/sys/net/ipv4/ip_forward`"
echo "done."
;;
*)
echo "Usage: $0 {start|restart|status}"
;;
esac
```

If you want to automatically have the network setup at boot time, you must create a softlink to `/etc/rc3.d/S99XenLaptopNAT`

When modifying the `modprobe.conf` file, you must include these lines:

```
alias dummy0 dummy
options dummy numdummies=1
```

18. Starting Domains Automatically During System Boot

Starting Domains Automatically During System Boot

You can configure your guests to start automatically when you boot the system. To do this, you

must modify the symbolic links that resides in `/etc/xen/auto` . This file points to the guest configuration files that you need to start automatically. The startup process is serialized, meaning that the higher the number of guests, the longer the boot process will take. This example shows you how to use symbolic links for the guest `rhel5vm01` :

```
[root@python xen]# cd /etc/xen
[root@python xen]# cd auto
[root@python auto]# ls
[root@python auto]# ln -s ../rhel5vm01 .
[root@python auto]# ls -l

lrwxrwxrwx 1 root root 14 Dec 14 10:02 rhel5vm01 -> ../rhel5vm01

[root@python auto]#
```

19. Modifying Domain0

To use Red Hat Virtualization to manage domain0, you will constantly making changes to the `grub.conf` configuration file, that resides in the `/etc` directory. Because of the large number of domains to manage, many system administrators prefer to use the 'cut and paste' method when editing `grub.conf` . If you do this, make sure that you include all five lines in the Virtualization entry (or this will create system errors). If you require Xen hypervisor specific values, you must add them to the 'xen' line. This example represents a correct `grub.conf` Virtualization entry:

```
# boot=/dev/sda/
default=0
timeout=15
#splashimage=(hd0, 0)/grub/splash.xpm.gz
hiddenmenu
serial --unit=0 --speed=115200 --word=8 --parity=no --stop=1
terminal --timeout=10 serial console
title Red Hat Enterprise Linux Server (2.6.17-1.2519.4.21. el5xen)
    root (hd0, 0)
    kernel /xen.gz-2.6.17-1.2519.4.21.el5 com1=115200, 8n1
    module /vmlinuz-2.6.17-1.2519.4.21el5xen ro
root=/dev/VolGroup00/LogVol00
    module /initrd-2.6.17-1.2519.4.21.el5xen.img
```

For example, if you need to change your dom0 hypervisor's memory to 256MB at boot time, you must edit the 'xen' line and append it with the correct entry, '`dom0_mem=256M`' . This example represents the respective `grub.conf` xen entry:

```
# boot=/dev/sda
default=0
timeout=15
#splashimage=(hd0,0)/grubs/splash.xpm.gz
hiddenmenu
```

```
serial --unit=0 --speed=115200 --word=8 --parity=no --stop=1
terminal --timeout=10 serial console
title Red Hat Enterprise Linux Server (2.6.17-1.2519.4.21.el5xen)
    root (hd0,0)
    kernel /xen.gz-2.6.17-1.2519.4.21.el5 com1=115200, 8n1
dom0_mem=256MB
    module /vmlinuz-2.6.17-1.2519.4.21.el5xen ro
    root=/dev/VolGroup00/LogVol00
    module /initrd-2.6.17-1.2519.4.21.el5xen.img
```

20. Guest Configuration Files

When you install new guests using `virt-manager` (or `virt-install`) tool(s) from Red Hat Enterprise Linux 5.1 with Virtualization, the guests configuration files (located in the `/etc/xen` directory) get modified and setup automatically. This configuration file example is for a para-virtualized guest:

```
name = "rhel5vm01"
memory = "2048"
disk = [ 'tap:aio:/xen/images/rhel5vm01.dsk,xvda,w', ]
vif = [ "type=ieomu, mac=00:16:3e:09:f0:12 bridge=xenbr0",
        "type=ieomu, mac=00:16:3e:09:f0:13 ]
vnc = 1
vncunused = 1
uuid = "302bd9ce-4f60-fc67-9e40-7a77d9b4e1ed"
bootloader = "/usr/bin/pygrub"
vcpus=2
on_reboot = "restart"
on_crash = "restart"
```

Note that the `serial="pty"` is the default for the configuration file. This configuration file example is for a fully-virtualized guest:

```
name = "rhel5u5-86_64"
builder = "hvm"
memory = 500
disk = [ 'file:/xen/images/rhel5u5-x86_64.dsk.hda,w
[../../../../home/mhideo/.evolution//xen/images/rhel5u5-x86_64.dsk.hda,w]' ]
vif = [ 'type=ioemu, mac=00:16:3e:09:f0:12, bridge=xenbr0', 'type=ieomu,
mac=00:16:3e:09:f0:13, bridge=xenbr1' ]
uuid = "b10372f9-91d7-ao5f-12ff-372100c99af5"
device_model = "/usr/lib64/xen/bin/qemu-dm"
kernel = "/usr/lib/xen/boot/hvmloader/"
vnc = 1
vncunused = 1
apic = 1
acpi = 1
pae = 1
vcpus = 1
serial = "pty" # enable serial console
on_boot = 'restart'
```

21. Cloning the Guest Configuration Files

You can copy (or clone) an existing configuration file to create an all new guest. You must modify the name parameter of the guests' configuration file. The new, unique name then appears in the hypervisor and is viewable by the management utilities. You must generate an all new UUID as well (using the `uuidgen(1)` command). Then for the `vif` entries you must define a unique MAC address for each guest (if you are copying a guest configuration from an existing guest, you can create a script to handle it). For the xen bridge information, if you move an existing guest configuration file to a new host, you must update the `xenbr` entry to match your local networking configuration. For the Device entries, you must modify the entries in the `'disk='` section to point to the correct guest image.

You must also modify these system configuration settings on your guest. You must modify the `HOSTNAME` entry of the `/etc/sysconfig/network` file to match the new guest's hostname.

You must modify the `HWADDR` address of the `/etc/sysconfig/network-scripts/ifcfg-eth0` file to match the output from `ifconfig eth0` file and if you use static IP addresses, you must modify the `IPADDR` entry.

22. Creating a Script to Generate MAC Addresses

Red Hat Virtualization can generate a MAC address for each virtual machine at the time of creation. Since is a nearly unlimited amount of numbers on the same subnet, it is unlikely you could get the same MAC address. To work around this, you can also write a script to generate a MAC address. This example script contains the parameters to generate a MAC address:

```
#!/usr/bin/python
# macgen.py script generates a MAC address for Xen guests
#
import random
mac = [ 0x00, 0x16, 0x3e,
        random.randint(0x00, 0x7f),
        random.randint(0x00, 0xff),
        random.randint(0x00, 0xff) ]
print ''.join(map(lambda x: "%02x" % x, mac))

Generates e.g.:
00:16:3e:66:f5:77
to stdout
```

23. Configuring Virtual Machine Live Migration

Red Hat Virtualization can migrate virtual machines between other servers running Red Hat Enterprise Linux 5.1 with Virtualization. Further, migration is performed in an offline method (using the `xm migrate` command). Live migration can be done from the same command.

However there are some additional modifications that you must do to the `xend-config` configuration file. This example identifies the entries that you must modify to ensure a successful migration:

`(xend-relocation-server yes)`

The default for this parameter is 'no', which keeps the relocation/migration server deactivated (unless on a trusted network) and the domain virtual memory is exchanged in raw form without encryption.

`(xend-relocation-port 8002)`

This parameter sets the port that `xend` uses for migration. This value is correct, just make sure to remove the comment that comes before it.

`(xend-relocation-address)`

This parameter is the address that listens for relocation socket connections, after you enable the `xend-relocation-server` . When listening, it restricts the migration to a particular interface.

`(xend-relocation-hosts-allow)`

This parameter controls the host that communicates with the relocation port. If the value is empty, then all incoming connections are allowed. You must change this to a space-separated sequences of regular expressions (such as `xend-relocation-hosts-allow- '^localhost\\.localdomain$'`). A host with a fully qualified domain name or IP address that matches these expressions are accepted.

After you configure these parameters, you must reboot the host for the Red Hat Virtualization to accept your new parameters.

24. Interpreting Error Messages

You receive the following error:

```
failed domain creation due to memory shortage, unable to balloon domain0
```

A domain can fail if there is not enough RAM available. Domain0 does not balloon down enough to provide space for the newly created guest. You can check the `xend.log` file for this error:

```
[2006-12-21] 20:33:31 xend 3198] DEBUG (balloon:133) Balloon: 558432 Kib
free; 0 to scrub; need 1048576; retries: 20
[2006-12-21] 20:33:31 xend. XendDomainInfo 3198] ERROR (XendDomainInfo: 202
Domain construction failed
```

You can check the amount of memory in use by domain0 by using the `xm list Domain0` command. If domain0 is not ballooned down, you can use the command `"xm mem-set`

Domain-0 NewMemSize" to check memory.

You receive the following error:

```
wrong kernel image: non-PAE kernel on a PAE
```

This message indicates that you are trying to run an unsupported guest kernel image on your Hypervisor. This happens when you try to boot a non-PAE paravirtual guest kernel on a RHEL 5.1 hypervisor. Red Hat Virtualization only supports guest kernels with PAE and 64bit architectures.

Type this command:

```
[root@smith]# xm create -c va base

Using config file "va-base"
Error: (22, 'invalid argument')
[2006-12-14 14:55:46 xend.XendDomainInfo 3874] ERRORs
(XendDomainInfo:202) Domain construction failed

Traceback (most recent call last)
File "/usr/lib/python2.4/site-packages/xen/xend/XendDomainInfo.py", line 195
in create vm.initDomain()
File " /usr/lib/python2.4/site-packages/xen/xend/XendDomainInfo.py", line
1363 in initDomain raise VmError(str(exn))
VmError: (22, 'Invalid argument')
[2006-12-14 14:55:46 xend.XendDomainInfo 3874] DEBUG (XenDomainInfo: 1449]
XendDlomainInfo.destroy: domin=1
[2006-12-14 14:55:46 xend.XendDomainInfo 3874] DEBUG (XenDomainInfo: 1457]
XendDlomainInfo.destroy:Domain(1)
```

If you need to run a 32bit/non-PAE kernel you will need to run your guest as a fully virtualized virtual machine. For paravirtualized guests, if you need to run a 32bit PAE guest, then you must have a 32bit PAE hypervisor. For paravirtualized guests, if you need to run a 64bit PAE guest, then you must have a 64bit PAE hypervisor. For full virtualization guests you must run a 64bit guest with a 64bit hypervisor. The 32bit PAE hypervisor that comes with RHEL 5 i686 only supports running 32bit PAE paravirtualized and 32 bit fully virtualized guest OSes. The 64bit hypervisor only supports 64bit paravirtualized guests.

This happens when you move the full virtualized HVM guest onto a RHEL 5.1 system. Your guest may fail to boot and you will see an error in the console screen. Check the PAE entry in your configuration file and ensure that pae=1. You should use a 32bit distribution.

You receive the following error:

```
Unable to open a connection to the Xen hypervisor or daemon
```

This happens when the virt-manager application fails to launch. This error occurs when there is no localhost entry in the `/etc/hosts` configuration file. Check the file and verify if the localhost entry is enabled. Here is an example of an incorrect localhost entry:

```
# Do not remove the following line, or various programs
# that require network functionality will fail.
localhost.localdomain localhost
```

Here is an example of a correct localhost entry:

```
# Do not remove the following line, or various programs
# that require network functionality will fail.
127.0.0.1 localhost.localdomain localhost
localhost.localdomain. localhost
```

You receive the following error (in the `xen-xend.log` file):

```
Bridge xenbr1 does not exist!
```

This happens when the guest's bridge is incorrectly configured and this forces the Xen hotplug scripts to timeout. If you move configuration files between hosts, you must ensure that you update the guest configuration files to reflect network topology and configuration modifications. When you attempt to start a guest that has an incorrect or non-existent Xen bridge configuration, you will receive the following errors:

```
[root@trumble virt]# xm create r5b2-mysql01

Using config file " r5b2-mysql01"
Going to boot Red Hat Enterprise Linux Server (2.6.18.-1.2747 .el5xen)
kernel: /vmlinuz-2.6.18-12747.el5xen
initrd: /initrd-2.6.18-1.2747.el5xen.img
Error: Device 0 (vif) could not be connected. Hotplug scripts not working.
```

In addition, the `xend.log` displays the following errors:

```
[2006-11-14 15:07:08 xend 3875] DEBUG (DevController:143) Waiting for
devices vif
[2006-11-14 15:07:08 xend 3875] DEBUG (DevController:149) Waiting for 0
[2006-11-14 15:07:08 xend 3875] DEBUG (DevController:464)
hotplugStatusCallback

/local/domain/0/backend/vif/2/0/hotplug-status
```

```
[2006-11-14 15:08:09 xend.XendDomainInfo 3875] DEBUG (XendDomainInfo:1449)
XendDomainInfo.destroy: domid=2
[2006-11-14 15:08:09 xend.XendDomainInfo 3875] DEBUG (XendDomainInfo:1457)
XendDomainInfo.destroyDomain(2)
[2006-11-14 15:07:08 xend 3875] DEBUG (DevController:464)
hotplugStatusCallback

/local/domain/0/backend/vif/2/0/hotplug-status
```

To resolve this problem, you must edit your guest configuration file, and modify the `vif` entry. When you locate the `vif` entry of the configuration file, assuming you are using `xenbr0` as the default bridge, ensure that the proper entry resembles the following:

```
# vif = ['mac=00:16:3e:49:1d:11, bridge=xenbr0',]
```

You receive these python deprecation errors:

```
[root@python xen]# xm shutdown win2k3xen12
[root@python xen]# xm create win2k3xen12

Using config file "win2k3xen12".

/usr/lib64/python2.4/site-packages/xenxm/opts.py:520: Deprecation Warning:
Non ASCII character '\xc0' in file win2k3xen12 on line 1, but no encoding
declared; see http://www.python.org/peps/pep-0263.html for details

execfile (defconfig, globs, locs,)
Error: invalid syntax 9win2k3xen12, line1)
```

Python generates these messages when an invalid (or incorrect) configuration file. To resolve this problem, you must modify the incorrect configuration file, or you can generate a new one.

25. Online Troubleshooting Resources

- Red Hat Virtualization Center

```
http://www.openvirtualization.com [http://www.openvirtualization.com/]
```

- Red Hat Enterprise Linux 5 Beta 2 Documentation

```
http://www.redhat.com/docs/manuals/enterprise/RHEL-5-manual/index.html
```

- Libvirt API

<http://www.libvirt.org> [http://www.libvirt.org/]

- virt-manager Project Home Page

<http://virt-manager.et.redhat.com> [http://virt-manager.et.redhat.com/]

- Xen Community Center

<http://www.xensource.com/xen/xen/>

- Virtualization Technologies Overview

<http://virt.kernelnewbies.org> [http://virt.kernelnewbies.org/]

- Emerging Technologies Projects

<http://et.redhat.com> [http://et.redhat.com/]

Additional Resources

To learn more about Red Hat Virtualization, refer to the following resources.

1. Useful Websites

- <http://www.cl.cam.ac.uk/research/srg/netos/xen/> — The project website of the Xen™ para-virtualization machine manager from which Red Hat Virtualization is derived. The site maintains the upstream Xen project binaries and sourcecode and also contains information, architecture overviews, documentation, and related links regarding Xen and its associated technologies.
- <http://www.libvirt.org/> — The official website for the `libvirt` virtualization API that interacts with the virtualization framework of a host OS.
- <http://virt-manager.et.redhat.com/> — The project website for the **Virtual Machine Manager** (`virt-manager`), the graphical application for managing virtual machines.

2. Installed Documentation

- `/usr/share/doc/xen-<version-number>/` — This directory contains a wealth of information about the Xen para-virtualization hypervisor and associated management tools, including a look at various example configurations, hardware-specific information, and the current Xen upstream user documentation.
- `man virsh` and `/usr/share/doc/libvirt-<version-number>` — Contains subcommands and options for the `virsh` virtual machine management utility as well as comprehensive information about the `libvirt` virtualization library API.
- `/usr/share/doc/gnome-applet-vm-<version-number>` — Documentation for the GNOME graphical panel applet that monitors and manages locally-running virtual machines.
- `/usr/share/doc/libvirt-python-<version-number>` — Provides details on the Python bindings for the `libvirt` library. The `libvirt-python` package allows python developers to create programs that interface with the `libvirt` virtualization management library.
- `/usr/share/doc/python-virtinst-<version-number>` — Provides documentation on the `virt-install` command that helps in starting installations of Fedora and Red Hat Enterprise Linux related distributions inside of virtual machines.
- `/usr/share/doc/virt-manager-<version-number>` — Provides documentation on the Virtual Machine Manager, which provides a graphical tool for administering virtual machines.

Appendix A. Lab 1

Xen Guest Installation

Goal: To install RHEL 3, 4, or 5 and Windows XP Xen guests.

Prerequisites: A workstation installed with Red Hat Enterprise Linux 5.0 with Virtualization component.

For this lab, you will configure and install RHEL 3, 4, or 5 and Win XP Xen guests using various virtualization tools.

Lab Sequence 1: Checking for PAE support

You must determine whether your system has PAE support. Red Hat Virtualization supports x86_64 or ia64 based CPU architectures to run para-virtualized guests. To run i386 guests the system requires a CPU with PAE extensions. Many older laptops (particularly those based on Pentium Mobile or Centrino) do not support PAE.

1. To determine if your CPU has PAE support, type:

```
grep pae /proc/cpuinfo
```

2. The following output shows a CPU that has PAE support. If the command returns nothing, then the CPU does not have PAE support. All the lab exercises require a i386 CPU with PAE extension or x86_64 or ia64 in order to proceed.

```
flags :  
fpu vme de pse tsc msr pae mce cx8 apic mtrr pge mca cmov pat clflush  
dts acpi  
mmx fxsr sse sse2 ss tm pbe nx up est tm2
```

Lab Sequence 2: Installing RHEL5 Beta 2 Xen para-virtualized guest using `virt-install`.

For this lab, you must install a Red Hat Enterprise Linux 5 Beta 2 Xen guest using `virt-install`.

1. To install your Red Hat Enterprise Linux 5 Beta 2 Xen guest, at the command prompt type:
`virt-install`.
2. When asked to install a fully virtualized guest, type: `no`.

3. Type `rhel5b2-pv1` for your virtual machine name.
4. Type 500 for your RAM allocation.
5. Type `/xen/rhel5b2-pv1.img` for your disk (guest image).
6. Type 6 for the size of your disk (guest image).
7. Type `yes` to enable graphics support.
8. Type `nfs:server:/path/to/rhel5b2` for your install location.
9. The installation begins. Proceed as normal with the installation.
- 10 After the installation completes, type `/etc/xen/rhel5b2-pv1`, and make the following changes: `#vnc=1#vncunused=1sdl=1`
- 11 Use a text editor to modify `/etc/inittab`, and append this to the file: `init`
`5.#id:3:initdefault:id:5:initdefault:`

Lab Sequence 3: Installing RHEL5 Beta 2 Xen para-virtualized guest using `virt-manager`.

For this lab, you will install a Red Hat Enterprise Linux 5 Beta 2 Xen paravirtualized guest using `virt-manager`.

1. To install your Red Hat Enterprise Linux 5 Beta 2 Xen guest, at the command prompt type:
`virt-manager`.
2. On the Open Connection window, select Local Xen host, and click on **Connect**.
3. Start Red Hat's Virtual Machine Manager application, and from the **File** menu, click on **New**.
4. Click on **Forward**.
5. Type `rhel5b2-pv2` for your system name, and click on **Forward**.
6. Select **Paravirtualized**, and click **Forward**.
7. Type `nfs:server:/path/to/rhel5b2` for your install media URL, and click **Forward**.
8. Select **Simple File**, type `/xen/rhel5b2-pv2.img` for your file location. Choose 6000 MB, and click **Forward**.
9. Choose 500 for your VM Startup and Maximum Memory, and click **Forward**.
- 10 Click **Finish**.

The Virtual Machine Console window appears. Proceed as normal and finish up the installation.

Lab Sequence 4: Checking for Intel-VT or AMD-V support

For this lab, you must determine if your system supports Intel-VT or AMD-V hardware. Your system must support Intel-VT or AMD-V enabled CPUs to successfully install the fully virtualized guest operating systems. Red Hat Virtualization incorporates a generic HVM layer to support these CPU vendors.

1. To determine if your CPU has Intel-VT or AMD-V support, type the following command:

```
egrep -e 'vmx|svm' /proc/cpuinfo
```

2. The following output shows a CPU that supports Intel-VT:

```
.flags :  
    fpu tsc msr pae mce cx8 apic mtrr mca cmov pat clflush dts acpi mmx fxsr  
sse  
    sse2 ss ht tm pbe constant_tsc pni monitor vmx est tm2 xtpr
```

If the command returns nothing, then the CPU does not support Intel-VT or AMD-V.

3. To determine if your CPU has Intel-VT or AMD-V support, type the following command:

```
cat /sys/hypervisor/properties/capabilities
```

4. The following output shows that Intel-VT support has been enabled in the BIOS. If the command returns nothing, then go into the BIOS Setup Utility and look for a setting related to 'Virtualization', i.e. 'Intel(R) Virtualization Technology' under 'CPU' section on a IBM T60p. Enable and save the setting and do a power off to take effect.

```
xen-3.0-x86_32p hvm-3.0-x86_32 hvm-3.0-x86_32p
```

Lab Sequence 5: Installing RHEL5 Beta 2 Xen fully virtualized guest using virt-install.

For this lab, you will install a Red Hat Enterprise Linux 5 Beta 2 Xen fully virtualized guest using virt-install:

1. To install your Red Hat Enterprise Linux 5 Beta 2 Xen guest, at the command prompt type:

```
virt-install.
```

2. When prompted to install a fully virtualized guest, type `yes`.

3. Type `rhel5b2-pv2` for your virtual machine name.

4. Type `500` for your memory allocation.

5. Type `/xen/rhel5b2-fv1.img` for your disk (guest image).
6. Type 6 for the size of your disk (guest image).
7. Type `yes` to enable graphics support.
8. Type `/dev/cdrom` for the virtual CD image.
9. The VNC viewer appears within the installation window. If there is an error message that says “main: Unable to connect to host: Connection refused (111)”, then type the following command to proceed: `vncviewer localhost:5900`. VNC port 5900 refers to the first Xen guest that is running on VNC. If it doesn't work, you might need to use 5901, 5902, etc.

The installation begins. Proceed as normal with the installation.

Lab Sequence 6: Installing RHEL5 Beta 2 Xen fully virtualized guest using `virt-manager`.

For this lab, you will install a Red Hat Enterprise Linux 5 Beta 2 Xen fully virtualized guest using `virt-manager`:

1. To install your Red Hat Enterprise Linux 5 Beta 2 Xen guest, at the command prompt type:
`virt-manager`.
2. On the **Open Connection** window, select Local Xen host, and click on **Connect**.
3. Start Red Hat's Virtual Machine Monitor application, and from the **File** menu, click on **New**.
4. Click on **Forward**.
5. Type `rhel5b2-fv2` for your system name, and click on **Forward**.
6. Select Fully virtualized, and click **Forward**.
7. Specify either CD-ROM or DVD, and enter the path to install media. Specify ISO Image location if you will install from an ISO image. Click **Forward**.
8. Select Simple File, type `/xen/rhel5b2-fv2.img` for your file location. Specify 6000 MB, and click **Forward**.
9. Choose 500 for your VM Startup and Maximum Memory, and click **Forward**.
10. Click **Finish**.
11. The Virtual Machine Console window appears.

Proceed as normal and finish up the installation.

Lab Sequence 7: Installing RHEL3 Xen fully virtualized guest using `virt-manager`.

For this lab, you will install a Red Hat Enterprise Linux 3 Xen guest using `virt-manager`:

1. The same instructions for Lab Sequence 6 applies here.

Lab Sequence 8: Installing RHEL4 Xen fully virtualized guest using `virt-manager`

For this lab, you will install a Red Hat Enterprise Linux 4 Xen guest using `virt-manager` :

1. The same instructions for Lab Sequence 6 applies here.

Lab Sequence 9: Installing Windows XP Xen fully virtualized guest using `virt-manager`.

For this lab, you will install a Windows XP Xen fully virtualized guest using `virt-manager`:

1. To install your Red Hat Enterprise Linux 5 on your Windows XP host, at the command prompt type: `virt-manager`.
2. On the **Open Connection** window, select Local Xen host, and click on **Connect**.
3. Start Red Hat's Virtual Machine Manager application, and from the **File** menu click on **New**.
4. Click on **Forward**.
5. Type `winxp` for your system name, and click on **Forward**.
6. Select Fully virtualized, and click **Forward**.
7. Specify either CD-ROM or DVD, and enter the path to install media. Specify ISO Image location if you will install from an ISO image. Click **Forward**.
8. Select Simple File, type `/xen/winxp.img` for your file location. Specify 6000 MB, and click **Forward**.
9. Select 1024 for your VM Startup and Maximum Memory, and select 2 for VCPUs. Click **Forward**.
10. Click **Finish**.
11. The Virtual Machine Console window appears. Proceed as normal and finish up the installation.
12. Choose to format the `C:\` partition in FAT file system format. Red Hat Enterprise Linux 5 does not come with NTFS kernel modules. Mounting or writing files to the Xen guest image may not be as straight-forward if you were to format the partition in NTFS file system format.
13. After you reboot the system for the first time, edit the `winxp` guest image:

```
losetup /dev/loop0 /xen/winxp.img  
kpartx -av /dev/loop0  
mount /dev/mapper/loop0p1 /mnt  
cp -prv $WINDOWS/i386 /mnt/. This fixes a problem that you may face in the later part of the Windows installation.
```
14. Restart the Xen guest manually by typing: `xm create -c winxp/`.

- 15 In the Virtual Machine Manager window, select the `winxp` Xen guest and click **Open**.
- 16 The Virtual Machine Console window appears. Proceed as normal and finish up with the installation.
- 17 Whenever a 'Files Needed' dialog box appears, change the path `GLOBALROOT\DEVICE\CDROM0\I386` to `C:\I386`. Depending on your installation, you may or may not see this problem. You may be prompted for missing files during the installation. Changing the path to `C:\I386` should compensate for this problem.
- 18 If the Xen guest console freezes, click `shutdown`, make the following changes in `/etc/xen/winxp: #vnc=1#vncunused=1sdl=1#vcpus=2`
- 19 Repeat step 14 and proceed as normal with the installation.

Appendix B. Lab 2

Live Migration

Goal: To configure and perform a live migration between two hosts.

Prerequisite: Two workstations installed with Red Hat Enterprise Linux 5.0 Beta 2 with Virtualization Platform, and a Fedora Core 6 Xen guest on one of the two workstations.

For this lab, you will configure the migration and execute a live migration between two hosts.

Introduction: Before you begin

For this lab, you will need two Virtualization hosts: a Xen guest and a shared storage. You must connect the two Virtualization hosts via a UTP cable. One of the Virtualization hosts exports a shared storage via NFS. You must configure both of the Virtualization hosts so they migrate successfully. The Xen guest resides on the shared storage. On the Xen guest, you should install a streaming server. You must make sure that the streaming server still runs without any interruptions on the Xen guest, so the live migration takes place between one Virtualization host and the other. For Lab 2, you will refer the two Virtualization hosts as `host1` and `host2`.

Sequence 1: Configuring xend (both Xen hosts)

In this Lab procedure, you configure xend to start up as a HTTP server and a relocation server. The xend daemon does not initiate the HTTP server by default. It starts the UNIX domain socket management server (for `xm`) and communicates with xend. To enable cross-machine live migration, you must configure it to support live migration:

1. To make a backup of your `xend-config.sxp` file:

```
cp -pr /etc/xen/xend-config.sxp /etc/xen/xend-config.sxp.default
```

2. Edit `/etc/xen/xend-config.sxp` and make the following changes:

```
 #(xend-unix-server yes)(xend-relocation-server
    yes)(xend-relocation-port 8002)(xend-relocation-address
    '')(xend-relocation-hosts-allow '')#(xend-relocation-hosts-allow
'^localhost$
    ^localhost\\.localdomain$')
```

3. Restart `xend:service` and `xend restart`.

Sequence 2: Exporting a shared storage via NFS

For this lab procedure, you will configure NFS and use it to export a shared storage.

1. Edit `/etc/exports` and include the line: `/xen *(rw,sync,no_root_squash)/`
2. Save `/etc/exports` and restart the NFS server. Make sure that the NFS server starts by default: `service nfs startchkconfig nfs on`.
3. After starting the NFS server on `host1`, we can then mount it on `host2`: `mount host1:/xen .`
4. Now start the Xen guest on `host1` and select `fc6-pv1` (or `fc6-pv2` from Lab 1):

```
xm create -c fc6-pv1
```

Sequence 3: Installing the Xen guest streaming server

For this lab step, you will install a streaming server, `gnump3d`, for our demonstration purposes. You will select `gnump3d` because it supports OGG vorbis files and is easy to install, configure, and modify.

1. Download `gnump3d-2.9.9.9.tar.bz2` tarball from <http://www.gnump3d.org/> . Unpack the tarball and in the `gnump3d-2.9.9.9/` directory, compile, and install the `gnump3d` application: `tar xvjf gnump3d-2.9.9.9.tar.bz2cd gnump3d-2.9.9.9/make install`
2. Create a `/home/mp3` directory and copy `TruthHappens.ogg` from Red Hat's Truth Happens page to `mkdir /home/mp3wget -c http://www.redhat.com/v/ogg/TruthHappens.ogg`
3. Start the streaming server by typing

```
command:gnump3d
```

4. On either one of the two Xen hosts, start running the Movie Player. If it is not installed, then install the `totem` and `iso-codecs` rpms before running the Movie Player. Click Applications, then Sound & Video, and finally Movie Player.

-
5. Click Movie, then Open Location. Enter `http://guest:8888/TruthHappens.ogg`.

Sequence 4: Performing live migration

1. Run the `TruthHappens.ogg` file on one of the two Xen hosts.
2. Perform the live migration from `host1` to `host2`:

```
xm migrate -live fc6-pv1 host2
```

3. Open multiple window terminals on both Xen hosts with the following command:

```
watch -n1 xm list
```

4. Observe as the live migration begins. Note how long it takes for migration to complete.

Challenge Sequence: Configuring VNC server from within the Xen guest

If time permits, from within the Xen guest, configure the VNC server to initiate when `gdm` starts up. Run VNC viewer and connect to the Xen guest. Play with the Xen guest when the live migration occurs. Attempt to pause/resume, and save/restore the Xen guest and observe what happens to the VNC viewer. If you connect to the VNC viewer via `localhost:590x`, and do a live migration, you won't be able to connect to the VNC viewer again when it dies. This is a known bug.

